



spring

Microservices with Spring

Lab Instructions

Building Cloud Native Applications using Spring
and Pivotal Cloud Foundry

Version 1.0

Copyright Notice

- Copyright © 2016 Pivotal Software, Inc. All rights reserved. This manual and its accompanying materials are protected by U.S. and international copyright and intellectual property laws.
- Pivotal products are covered by one or more patents listed at <http://www.pivotal.io/patents>.
- Pivotal is a registered trademark or trademark of Pivotal Software, Inc. in the United States and/or other jurisdictions. All other marks and names mentioned herein may be trademarks of their respective companies. The training material is provided “as is,” and all express or implied conditions, representations, and warranties, including any implied warranty of merchantability, fitness for a particular purpose or noninfringement, are disclaimed, even if Pivotal Software, Inc., has been advised of the possibility of such claims. This training material is designed to support an instructor-led training course and is intended to be used for reference purposes in conjunction with the instructor-led training course. The training material is not a standalone training tool. Use of the training material for self-study without class attendance is not recommended.
- These materials and the computer programs to which it relates are the property of, and embody trade secrets and confidential information proprietary to, Pivotal Software, Inc., and may not be reproduced, copied, disclosed, transferred, adapted or modified without the express written approval of Pivotal Software, Inc.

This Page Intentionally Left Blank

This Page Intentionally Left Blank

Pivotal

Table of Contents

1. Course Setup	1
1.1. Note on Lab Usage	1
1.2. Java	1
1.3. IDE	2
1.4. Maven	2
1.5. Git Command-Line Utility	3
1.6. Cloud Foundry Command-Line Utility.....	3
1.7. Curl for Windows	3
1.8. Optional Add-on - View JSON Data.....	3
1.8.1. Firefox or Chrome.....	3
1.8.2. Internet Explorer	4
1.9. Testing Your Setup.....	4
2. Lab 1a - Getting Started with Spring Boot.....	5
2.1. Generate a project from start.spring.io	5
2.2. Run it locally	5
3. Lab 1b - Deploying a Web Application with an Embedded Container	7
3.1. Build and Run with Embedded Apache Tomcat	7
3.2. Hot reload with Spring Boot Devtools.....	7
3.3. Build and Run with Embedded Eclipse Jetty	8
4. Lab 1c - Externalizing Configuration with Spring Boot	10
4.1. Refactoring to Externalize the Config	10
4.2. Using Environment Variables for Config	10
5. Lab 1d - Introspection, Monitoring, and Metrics using Spring Boot Actuator.....	12
5.1. Set up the Actuator	12
5.2. Adding proper JSON formatting	12
5.3. Introspection Endpoints.....	12
5.4. Health Indicators.....	13
6. Lab 2a - From Zero to Pushing Your First Application.....	15
6.1. Setup	15
6.2. Build and Push!	15
7. Lab 2b - Binding to Cloud Foundry Services.....	18
7.1. A Bit of Review.....	18
7.2. The Services Marketplace	19
7.3. Creating and Binding to a Service Instance	21
7.4. Optional: Swapping from MySQL to MongoDB	23
7.5. Clean Up.....	24

8. Lab 2c - Scaling Applications	26
8.1. Push the cf-scale-boot Application.....	26
8.2. Scale the Application Up	27
8.2.1. About the Router	28
8.3. Scale the Application Down	29
9. Lab 2d - Monitoring Applications	31
9.1. Events	31
9.2. Logs	32
9.2.1. For Developers	33
9.2.2. For Operators (OPTIONAL).....	34
9.3. Health	36
9.4. Clean Up.....	39
10. Lab 03a - Build a Hypermedia-Driven RESTful Web Service with Spring Data REST	40
10.1. Initializing the Application	40
10.2. Importing Data	42
10.3. Adding Search.....	43
10.4. Pushing to Cloud Foundry	44
11. Lab 03b - Leveraging Spring Cloud Connectors for Service Binding	46
11.1. Using Spring Cloud Connectors	46
11.2. Customising the DataSource	48
12. Lab 04a - Identify the bounded contexts of the Monolith App	49
12.1. Understanding the spring-trader monolith application.....	49
12.2. Decomposition into microservices.....	50
13. Lab 05a - Build a Quotes services using MongoDB.....	51
13.1. Exploring springtrader-quotes.....	51
13.2. Preparing for Cloud Foundry	52
13.3. Deploying to Cloud Foundry	53
14. Lab 05b - Build Accounts Service with Mysql/H2	55
14.1. Exploring springtrader-accounts	55
14.2. Preparing for Cloud Foundry	56
14.3. Deploying to Cloud Foundry	57
15. Lab 05c - Build a Portfolio with MySQL	58
15.1. Exploring springtrader-portfolio	58
15.2. Preparing for Cloud Foundry	59
15.3. Deploying to Cloud Foundry	60
16. Lab 05d - Run the Web User Interface.....	62
16.1. Exploring springtrader-portfolio	62
16.2. Running the Application	63
16.3. Deploying to Cloud Foundry	63
17. Lab 06a - Deploying and Using Spring Cloud Config Server.....	65
17.1. Create the config folder to Store Configuration	65
17.2. Create a Spring Cloud Config Server	66

17.3. Create the Sample Test Application	67
17.4. [Bonus] Use github to store the config files.....	69
18. Lab 06b - Leveraging Eureka for Service Discovery via Spring Cloud Netflix	70
18.1. Introduction	70
18.2. Creating a Eureka Server.....	70
18.3. Create and Register the Producer Service	72
18.4. Create and Register the Consumer Service	73
19. Lab 07a - Client-Side Load Balancing with Ribbon	76
19.1. Setup.....	76
19.2. Using the LoadBalancerClient	77
19.3. Failing one of the producers.....	78
20. Lab 07b - Fault-Tolerance with Hystrix	79
20.1. Setup.....	79
20.2. Using Hystrix	80
20.3. Failover.....	81
20.4. Fallback	81
20.5. Recovery	82
20.6. Next Steps.....	82
21. Lab 07c - Monitoring Circuit Breakers with Hystrix Dashboard.....	83
21.1. Setup.....	83
21.2. Building the Hystrix Dashboard	83
21.3. Fallback	85
21.4. Recovery	86
22. Lab 07d - Declarative REST Clients with Feign.....	87
22.1. Setup	87
22.2. Using Feign	87
23. Lab 08a - Creating an OAuth2 Authorization Server.....	89
23.1. Default behaviour	89
23.2. Configuring the dependencies.....	89
23.3. Enabling the auth server.....	90
23.4. Testing your Oauth server	90
24. Lab 08b - Securing a Resource Server with Spring Cloud Security.....	92
24.1. Running it	92
24.2. Securing it.....	92
24.3. Obtain a valid token	93
25. Lab 08c - Consuming a Secured Resource from OAuth2-Aware Client	95
25.1. Create a OAuth2-Aware Client	95
25.2. Test your client.....	96
26. (Optional) Lab 09a - Deploying the StringTrader Microservices version Application	97
26.1. Copying the 3 services	97
26.2. Starting Spring Cloud Services on Pivotal Cloud Foundry	97
26.3. Deploying the Application	97

26.4. Testing it out.....	98
A. Eclipse Tips	102
A.1. Introduction.....	102
A.2. Package Explorer View	102
A.3. Add Unimplemented Methods	104
A.4. Field Auto-Completion	104
A.5. Generating Constructors From Fields.....	105
A.6. Field Naming Conventions	105
A.7. Tasks View	106
A.8. Rename a File	106

Chapter 1. Course Setup

The following software is required to complete this course:

- Java JDK 8
- An IDE or Editor of your choice
- Maven utility: `mvn`
- Git command line: `git`
- Cloud Foundry CLI `cf`
- Curl (Windows users only)
- Optional: Browser plugin to view JSON Data (Chrome or Firefox only)

Ideally this setup should be completed *before* you attend the course.

1.1. Note on Lab Usage

Much of this course involves using command line tools rather than GUI applications. This may be unfamiliar to some of you.

In particular, Windows users will need to open and use Command windows. To do this press WK + R and type `cmd` (where WK is the Windows Key).

To keep the labs simple, although we encourage you to use an IDE, we only use it for creating and editing files. In most cases we build using `mvn` directly. We accept that this is not typical but it makes the lab instructions simpler.

1.2. Java

A Java 8 JDK needs to be installed. Go to Oracle at <http://www.oracle.com/technetwork/java/javase/downloads/index.html> and click the Java download button on the right. You should be on the page to download the "Java SE Development Kit".

Click the radio-button to say you "Accept License Agreement" and then pick the right JDK for your platform. Make sure you are downloading the JDK *not* the JRE.

Once downloaded, run the installer. It takes a while to run.

1.3. IDE

We recommend installing the integrated development environment of your choice, such as Eclipse, Spring Tool Suite (STS) or IntelliJ. If you have no preference we recommend STS as it has additional support for the Spring facilities used during the course.

- STS: <http://spring.io/tools> The download is a zip file, so unpack it somewhere convenient. Windows users are recommended to unpack it directly into c:\ otherwise you will have problems exceeding the maximum path length (yes even on Windows 8 and 10).
Once unpacked you should find a directory similar to `sts_bundle/sts-X.X.X.RELEASE` or `sts_bundle/sts-X.X.X.RELEASE-eX.X.X` (the versions will differ). In here is the executable for STS (`STS.exe`, `STS.app` or `STS`). For convenience, you might like to create a shortcut on your Desktop.
- Eclipse: <https://eclipse.org/downloads/> Either use the Installer or choose "Eclipse IDE for Java Developers".
- IntelliJ: <https://www.jetbrains.com/idea/download/> The Community Edition is free and will be sufficient for this course.

An IDE is not required to do this course, but if you prefer not to use an IDE, you will need to use your favorite editor instead. Windows users are recommended to download the free Notepad++s editor from <https://notepad-plus-plus.org/download>

1.4. Maven

Download Maven from <https://maven.apache.org/download.cgi> - select the binary zip archive (Windows) or the binary tar.gz archive (MacOS, Linux).

Once downloaded, unpack into a convenient directory. You need to add the `bin` directory to your execution path.

- Windows - Right click `My Computer/This PC # Properties # Advanced system settings # Environment variables ...PATH` should already exist, click `Edit` and add to the end. Don't forget to add a semicolon first. For more details on how to do this see <http://www.computerhope.com/issues/ch000549.htm>.
- MacOS, Linux - Assuming you are using a BASH shell, edit `~/.bashrc` and add:

```
PATH=$PATH:/path/to/maven/bin
```

You will need to close any existing Terminal/Cmd windows and open a new one for the change to take effect.

1.5. Git Command-Line Utility

Go to <https://git-scm.com/downloads>, click the image corresponding to your platform, download the installer and run it.

Warning: Ignore the large computer monitor labelled "Latest source Release".

- Windows users: When running the installer, select "*Use Git from the Windows Command Prompt*" (this is not the default). Accept all other defaults.

1.6. Cloud Foundry Command-Line Utility

Go to <https://github.com/cloudfoundry/cli/releases>, download and run the right installer for your platform.

1.7. Curl for Windows

Linux and MacOS users already have the `curl` utility. Windows users can download it from here: <https://github.com/S2EDU/PCFIImmersionStudentFiles/raw/master/curl.zip>.

Once downloaded, unpack it somewhere convenient and add to your path - the same procedure as for Maven.

1.8. Optional Add-on - View JSON Data

Several labs use `curl` to send requests that return JSON data. If you prefer to do this in your browser ...

1.8.1. Firefox or Chrome

1. Do an Internet search for the "JSONView" add-on.
2. Install as directed
3. Once installed, try this JSON test page: <http://jsonplaceholder.typicode.com/posts/1>

1.8.2. Internet Explorer

1. Create a file on your desktop called `json.reg`.
2. Right click and select edit.
3. Copy and paste the following:

```
Windows Registry Editor Version 5.00;
; Tell IE 7,8,9,10,11 to open JSON documents in the browser on Windows XP and later.
; 25336920-03F9-11cf-8FD0-00AA00686F13 is the CLSID for the "Browse in place" .
;
[HKEY_CLASSES_ROOT\ MIME\ Database\ Content Type\ application/json]
"CLSID"="{25336920-03F9-11cf-8FD0-00AA00686F13}"
"Encoding"=hex:08,00,00,00
```

4. Save and close.
5. Double click on the `json.reg` to run it. You will get a warning, click `Yes` to continue.
6. Now try this JSON test page: <http://jsonplaceholder.typicode.com/posts/1>
7. Delete `json.reg`

1.9. Testing Your Setup

If everything has installed correctly, *open a brand new Terminal/Cmd window*. You should be able to run the following:

```
java -version
mvn -version
git --version
cf --version
curl --version
```

Chapter 2. Lab 1a - Getting Started with Spring Boot

Estimated time to complete: 20 minutes

2.1. Generate a project from start.spring.io

1. In your browser, visit <https://start.spring.io>.
2. Fill out the **Project metadata** fields as follows:

Group

io.spring.pivotal

Artifact

hello-spring-boot

Search For Dependencies

Web

Spring Boot version

1.3.5

3. Click the **Generate Project** button. Your browser will download a zip file. Unpack that zip file at \$COURSE_HOME/lab_01/lab_01a/initial.

2.2. Run it locally

1. Import the project's `pom.xml` into your editor/IDE of choice.
2. Add a `@RestController` annotation to the class `io.pivotalspring.hello.HelloSpringBootApplication` to indicate that we wish to return REST (data) responses not HTML.
It will be marked in red, so you need to add the import for this annotation.

Note



Your IDE can do this for you using the "[big red]*Organise imports" option. In Eclipse/STS

press CTRL-SHIFT-O (MacOS: COMMAND-SHIFT-O). In IntelliJ use CTRL+ALT+o (MacOS: COMMAND-ALT-o)

1. Add the following request handler method to the class `io.pivotal.spring.hello.HelloSpringBootApplication` after the static `main` method:

```
@RequestMapping(*/*)
public String hello() {
    return "Hello World!";
}
```

Again you will need to use the IDE to import the `@RequestMapping` annotation.

2. To run the application:

- Eclipse/STS: Right click in the editor window and select `Run As # Spring Boot App`.
- IntelliJ: In the editor window press CTRL-SHIFT-R (MacOS: COMMAND-SHIFT-R)
Wait for the application to start up.

3. Open <http://localhost:8080> in your browser and it should return `Hello World!`.

Congratulations! You've just completed your first Spring Boot application.

Chapter 3. Lab 1b - Deploying a Web Application with an Embedded Container

Estimated time to complete: 15 minutes

3.1. Build and Run with Embedded Apache Tomcat

Spring Boot embeds Apache Tomcat by default.



Note

For this lab, you should keep working with the project that you have created in Lab 01a

- Run the application by launching the `HelloSpringBootApplication` class from your IDE
You should see the application start up an embedded Apache Tomcat server on port 8080:

```
s.b.c.e.t.TomcatEmbeddedServletContainer : Tomcat started on port(s): 8080 (http)
i.p.s.hello.HelloSpringBootApplication   : Started HelloSpringBootApplication in 3.023 seconds (JVM running for 3.432)
```

- Visit the application in the browser (<http://localhost:8080>), and you should see the following:

```
Hello World!
```

3.2. Hot reload with Spring Boot Devtools

One of the coolest features of Spring Boot is that most of your code can hot redeploy (you do not have to restart your application for changes to be taken into account).

- Open your `pom.xml` file and add the following dependency:

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-devtools</artifactId>
</dependency>
```

- Restart your application once
- Inside `HelloSpringBootApplication`, replace the String "Hello World" with "Hello Wooooorld" and save the file. You will see that the application redeploys the changes automatically. You just have to go to your browser and refresh so you can see the new value.



Note

IntelliJ users, the editor auto saves your changes, however devtools only monitors the compiled .class files, so please use Build # Make Project



Note

You still need to wait a few seconds for Spring Boot to refresh your application but it takes 20-30% of the time that it would take to restart your application completely.

3.3. Build and Run with Embedded Eclipse Jetty

Spring Boot also supports embedding a Jetty server.

1. Open pom.xml and replace the following:

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-web</artifactId>
</dependency>
```

with:

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-web</artifactId>
  <exclusions>
    <exclusion>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-starter-tomcat</artifactId>
    </exclusion>
  </exclusions>
</dependency>
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-jetty</artifactId>
</dependency>
```

2. Build and run the application

You can run it within your Development environment or from the command line (same as earlier in this lab).

You should see the application start up an embedded Jetty server on port 8080:

```
s.b.c.e.j.JettyEmbeddedServletContainer : Jetty started on port(s) 8080 (http/1.1)
i.p.s.hello.HelloSpringBootApplication  : Started HelloSpringBootApplication in 3.671 seconds (JVM running for 4.079)
```

3. Visit the application in the browser (<http://localhost:8080>), and you should see the following:

```
Hello World!
```



Note

Hot redeploy works with Jetty in the same way as it works for Tomcat

Chapter 4. Lab 1c - Externalizing Configuration with Spring Boot

Estimated time to complete: 20 minutes

4.1. Refactoring to Externalize the Config

1. You should work using the same project as in the previous labs
2. Inside the folder `src/main/resources/`, you will see a file called `application.properties`. Rename that file into `application.yml`



Note

Spring Boot allows to work with properties files or YML files indifferently. In this workshop you will be using YML files because their syntax is more powerful (and still fairly simple!)

1. Inside `application.yml` paste the following:

```
greeting: Hello
```

2. To the class `io.pivotal.spring.hello.HelloSpringBootApplication`, add a `greeting` field and inject its value:

```
@Value("${greeting}")
String greeting = "Bonjour";
```

3. Also `io.pivotal.spring.hello.HelloSpringBootApplication`, change the return statement of `hello()` to the following:

```
return String.format("%s World!", greeting);
```

4. Make sure your application is running and visit <http://localhost:8080> in your browser. Which of the values is picked up? Is it the local one ("Bonjour") or the one in `application.yml` ("Hello")?

4.2. Using Environment Variables for Config

1. You should now stop the application inside your IDE because you are going to launch it from the command line

2. You are now going to add an environment variable from the command line. You should compile and run the application as shown below:

Windows.

```
> mvn clean package  
> set GREETING=Ohai && java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

Linux, MacOS:

```
$ mvn clean package  
$ GREETING=Ohai java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

 **Note**

in case of any issue running your application, you should confirm that your jar file is the same as in the above code (it varies depending on the name of our project)

1. Visit the application in the browser (<http://localhost:8080>), and verify that the output has changed to the following:

```
Ohai World!
```

2. Stop the application.
3. Visit <http://docs.spring.io/spring-boot/docs/current/reference/html/boot-features-external-config.html> to learn more about this outcome and the entire priority scheme for conflict resolution.

Chapter 5. Lab 1d - Introspection, Monitoring, and Metrics using Spring Boot Actuator

Estimated time to complete: 20 minutes

5.1. Set up the Actuator

1. You should use the same project as in the previous labs.
2. To pom.xml add the following dependency to include the starter for Spring Boot Actuator:

```
<dependency>
<groupId>org.springframework.boot</groupId>
<artifactId>spring-boot-starter-actuator</artifactId>
</dependency>
```
3. Run the application from your IDE
4. open this URL in your browser: <http://localhost:8080/beans> (use right-click Open in New Tab). You can see runtime information about your Spring beans. However it is hard to read.

5.2. Adding proper JSON formatting

1. Add the following configuration *at the end* of the file HelloSpringBootApplication.java:
(you are allowed to have more than one classes in the same file as long as only one of them is public)

```
@Configuration
class RestWsApplicationConfig { // Note that this class is NOT public

    @Bean
    public Jackson2ObjectMapperBuilder jacksonBuilder() {
        Jackson2ObjectMapperBuilder builder = new Jackson2ObjectMapperBuilder();
        builder.indentOutput(true);
        return builder;
    }
}
```

2. Run the application again and connect to <http://localhost:8080/beans>

5.3. Introspection Endpoints

1. Try out the following endpoints. The output is omitted here because it can be quite large:

<http://localhost:8080/beans>

Dumps all of the beans in the Spring context.

<http://localhost:8080/autoconfig>

Dumps all of the auto-configuration performed as part of application bootstrapping.

<http://localhost:8080/env>

Dumps the application's shell environment as well as all Java system properties.

<http://localhost:8080/mappings>

Dumps all URI request mappings and the controller methods to which they are mapped.

<http://localhost:8080/dump>

Performs a thread dump.

<http://localhost:8080/trace>

Displays trace information (by default the last few HTTP requests).

5.4. Health Indicators

Spring Boot provides an endpoint (<http://localhost:8080/health>) that allows for the notion of various health indicators. Try it now.

1. In production it would be normal to enable Spring Security so that the `/health` endpoint will only expose an `UP` or `DOWN` value unless you are an authorized user. However, to simplify working with the endpoint for this lab, we will leave security off (the default behavior).

You can experiment with this now by modifying the `sensitivity` property which is `false` by default. Add the following to `src/main/resources/application.yml` and refresh the health page:

```
endpoints:  
  health:  
    sensitive: true
```

Make sure to set the property back to `false` before continuing.

2. Create a *new* class `io.pivotal.spring.FlappingHealthIndicator` (in its *own* file) and into it paste the following code:

```
@Component  
public class FlappingHealthIndicator implements HealthIndicator{
```

Lab 1d - Introspection, Monitoring, and Metrics using Spring Boot Actuator

```
private Random random = new Random(System.currentTimeMillis());  
  
@Override  
public Health health() {  
    int result = random.nextInt(100);  
    if (result < 50) {  
        return Health.down().withDetail("flapper", "failure").withDetail("random", result).build();  
    } else {  
        return Health.up().withDetail("flapper", "ok").withDetail("random", result).build();  
    }  
}
```

This demo health indicator will randomize the health check.

3. Run the application
4. Visit the application in the browser (<http://localhost:8080/health>), and verify that the output is similar to the following (and changes randomly!):

```
{  
    "status": "UP",  
    "flapping": {  
        "status": "UP",  
        "flapper": "ok",  
        "random": 69  
    },  
    "diskSpace": {  
        "status": "UP",  
        "free": 113632186368,  
        "threshold": 10485760  
    }  
}
```

To learn more about the autogenerated metrics, visit <http://docs.spring.io/spring-boot/docs/current/reference/html/production-ready-metrics.html>.

Chapter 6. Lab 2a - From Zero to Pushing Your First Application

6.1. Setup

1. Set the API target for the CLI:

```
$ cf api {CF-API-URL} --skip-ssl-validation
```

2. Your username will be **workshop.user**. Password will be provided in the class.

3. Follow the steps above for login.

4. Create your own space and target that space:

```
$ cf create-space YOURNAME  
$ cf target -s YOURNAME
```

6.2. Build and Push!

1. Change to the *Spring Music* sample application directory:

```
$ cd $COURSE_HOME/lab/lab_02/lab_02a/spring-music
```

2. Using the Gradle Wrapper, build and package the application:

```
$ ./gradlew assemble
```

The Gradle Wrapper will automatically download the appropriate version of Gradle for this project along with all of *Spring Music*'s dependencies. This may take a few moments.

3. Push the application!

```
$ cf push
```

You should see output similar to the following listing. Take a look at the listing callouts for a play-by-play of what's happening:

```
Using manifest file .../lab_02/lab_02a/spring-music/manifest.yml  
❶  
Creating app spring-music in org oreilly-class / space instructor as example@pivotal.io...  
OK  
❷  
Creating route spring-music-hippest-shaman.cfapps.io...  
OK
```

```
①
Binding spring-music-hippest-shaman.cfapps.io to spring-music...
OK
②

Uploading spring-music...
③
Uploading app files from: .../lab_02/lab_02a/spring-music/build/libs/spring-music.war
Uploading 569.7K, 90 files
Done uploading
OK

Starting app spring-music in org oreilly-class / space instructor as example@pivotal.io...
④
----> Downloaded app package (21M)
----> Java Buildpack Version: v2.6.1 | https://github.com/cloudfoundry/java-buildpack.git#2d92e70
----> Downloading Open Jdk JRE 1.8.0_31 from https://download.run.pivotal.io/openjdk/x86_64/openjdk
     -1.8.0_31.tar.gz (1.3s)
      Expanding Open Jdk JRE to .java-buildpack/open_jdk_jre (1.2s)
⑤
----> Downloading Spring Auto Reconfiguration 1.7.0_RELEASE from https://download.run.pivotal.io/auto-reconfiguration
     /auto-reconfiguration-1.7.0_RELEASE.jar (0.1s)
----> Downloading Tomcat Instance 8.0.18 from https://download.run.pivotal.io/tomcat/tomcat-8.0.18.tar.gz (0.4s)
      Expanding Tomcat to .java-buildpack/tomcat (0.1s)
⑥
----> Downloading Tomcat Lifecycle Support 2.4.0_RELEASE from https://download.run.pivotal.io/tomcat-lifecycle-support
     /tomcat-lifecycle-support-2.4.0_RELEASE.jar (0.0s)
----> Downloading Tomcat Logging Support 2.4.0_RELEASE from https://download.run.pivotal.io/tomcat-logging-support
     /tomcat-logging-support-2.4.0_RELEASE.jar (0.0s)
----> Downloading Tomcat Access Logging Support 2.4.0_RELEASE from https://download.run.pivotal.io/tomcat-access
     -logging-support/tomcat-access-logging-support-2.4.0_RELEASE.jar (0.0s)
----> Uploading droplet (66M)
⑦
0 of 1 instances running, 1 starting
1 of 1 instances running

App started

OK

App spring-music was started using this command `JAVA_HOME=$PWD/.java-buildpack/open_jdk_jre JAVA_OPTS=*
-Djava.io.tmpdir=$TMPDIR
-XX:OnOutOfMemoryError=$PWD/.java-buildpack/open_jdk_jre/bin/killjava.sh -Xmx382293K -Xms382293K -XX:MaxMetaspaceSize=64
-M
-XX:MetaspaceSize=64M -Xss995K -Daccess.logging.enabled=false -Dhttp.port=$PORT` $PWD/.java-buildpack/tomcat/bin
/catalina.sh run`
⑧

Showing health and status for app spring-music in org oreilly-class / space instructor as example@pivotal.io...
OK

requested state: started
instances: 1/1
usage: 512M x 1 instances
urls: spring-music-hippest-shaman.cfapps.io
last uploaded: Fri Feb 13 15:43:08 UTC 2015

#0   state      since            cpu    memory         disk
#0   running   2015-02-13 09:43:55 AM  0.0%   394.5M of 512M  131.1M of 1G
```

- ① The CLI is using a manifest to provide necessary configuration details such as application name, memory to be allocated, and path to the application artifact. Take a look at `manifest.yml` to see how.
- ② In most cases, the CLI indicates each Cloud Foundry API call as it happens. In this case, the CLI has created an application record for *Spring Music* in your assigned space.
- ③ All HTTP/HTTPS requests to applications will flow through Cloud Foundry's front-end router called [\(Go\)Router](#). Here the CLI is creating a route with random word tokens inserted (again, see `manifest.yml` for a hint!) to prevent route collisions across the default `cfapps.io` domain.
- ④ Now the CLI is *binding* the created route to the application. Routes can actually be bound to multiple

applications to support techniques such as [blue-green deployments](#).

- ⑤ The CLI finally uploads the application bits to PWS. Notice that it's uploading *90 files!* This is because Cloud Foundry actually explodes a ZIP artifact before uploading it for caching purposes.
 - ⑥ Now we begin the staging process. The [Java Buildpack](#) is responsible for assembling the runtime components necessary to run the application.
 - ⑦ Here we see the version of the JRE that has been chosen and installed.
 - ⑧ And here we see the version of Tomcat that has been chosen and installed.
 - ⑨ The complete package of your application and all of its necessary runtime components is called a *droplet*. Here the droplet is being uploaded to PWS's internal blobstore so that it can be easily copied to one or more Execution Agents (dedicated VMs that run application instances).
 - ⑩ The CLI tells you exactly what command and argument set was used to start your application.
- Finally the CLI reports the current status of your application's health. You can get the same output at any time by typing `cf app spring-music`.

4. Visit the application in your browser using the route that was generated by the CLI:

The screenshot shows a web application titled "Spring Music" with a green header bar. Below the header, there is a navigation bar with the title "Albums". On the right side of the header, there is a status box showing "Profiles: cloud,in-memory" and "Services:". Below the header, there is a search bar with the placeholder "[view as: list | sort by: title artist year genre | +add an album]". The main content area displays a grid of 12 album cards, each containing the album title, artist, year, genre, and a small gear icon for more options. The albums listed are:

Album	Artist	Year	Genre
IV	Led Zeppelin	1971	Rock
Nevermind	Nirvana	1991	Rock
What's Going On	Marvin Gaye	1971	Rock
Are You Experienced?	Jimi Hendrix Experience	1967	Rock
The Joshua Tree	U2	1987	Rock
Abbey Road	The Beatles	1969	Rock
Rumours	Fleetwood Mac	1977	Rock
Sun Sessions	Elvis Presley	1976	Rock
Thriller	Michael Jackson	1982	Pop
Exile on Main Street	The Rolling Stones	1972	Rock
Born to Run	Bruce Springsteen	1975	Rock
London Calling	The Clash	1980	Rock

Be sure to click on the (i) *information icon* in the top right-hand corner of the UI. The drop-down gives you important information about the state of the currently running *Spring Music* instance, including what Spring Profiles are turned on and what Cloud Foundry services are bound. It will be important in the next lab!

Chapter 7. Lab 2b - Binding to Cloud Foundry Services

The *Spring Music* application was designed to illustrate the ease with which various types of data services can be bound to and utilized by Spring applications running on Cloud Foundry. In this lab, we'll be binding the application to both MySql and MongoDB databases.

Cloud Foundry services are managed through two primary types of operations:

Create/Delete

These operations create or delete instances of a service. For a database this could mean creating/deleting a schema in an existing multitenant cluster or creating/deleting a dedicated database cluster.

Bind/Unbind

These operations create or delete unique credential sets for an existing service instance that can then be injected into the environment of an application instance.

7.1. A Bit of Review

Your instance of *Spring Music* should still be running from the end of the [previous lab](#).

Visit the application in your browser by hitting the route that was generated by the CLI:

Spring Music ♪

Albums

[view as: | sort by: title artist year genre | [+add an album](#)]

IV Led Zeppelin 1971 Rock 	Nevermind Nirvana 1991 Rock 	What's Going On Marvin Gaye 1971 Rock 	Are You Experienced? Jimi Hendrix Experience 1967 Rock
The Joshua Tree U2 1987 Rock 	Abbey Road The Beatles 1969 Rock 	Rumours Fleetwood Mac 1977 Rock 	Sun Sessions Elvis Presley 1976 Rock
Thriller Michael Jackson 1982 Pop 	Exile on Main Street The Rolling Stones 1972 Rock 	Born to Run Bruce Springsteen 1975 Rock 	London Calling The Clash 1980 Rock

Profiles: cloud.in-memory
Services:

The information dialog in the top right-hand corner indicates that we're currently running with an in-memory database, and that we're not bound to any services. Let's change that.

7.2. The Services Marketplace

There are two ways to discover what services are available when using Pivotal Cloud Foundry (PCF):

1. The first is available on *any* instance of Cloud Foundry (Open Source or PCF) via the cf CLI. Type:

```
$ cf marketplace
```

and you'll get a list of services, their available plans, and descriptions (the services list may vary from what we show here):

```
instructor$ cf marketplace
Getting services from marketplace in org pivotaledu / space development as pchapman@gopivotal.com...
OK

service      plans                                         description
3scale       free_appdirect, basic_appdirect*, pro_appdirect*
app-autoscaler bronze, gold                                API Management Platform
                   (beta)                                     Scales bound applications in response to load
blazemeter   free-tier, basic1kmr*, pro5kmr*
cleardb      spark, boost*, amp*, shock*                      Performance Testing Platform
...                                     Highly available MySQL for your Apps.
```

2. The second way is specific to PCF's Application Manager UI. If using Pivotal Web Services (PWS), this is what you see once you login to <http://run.pivotal.io>. If you haven't already, login now.

Click on the “Marketplace” link (on the left hand side panel):*The Screenshots below are from PWS (public managed Pivotal CF). If you are using your own installation of PCF, you will see something similar.*

The screenshot shows the Pivotal Web Services Application Manager interface. On the left, there's a sidebar with 'ORG' set to 'oreilly-class' and 'SPACES' listed: 'development', 'instructor' (which is selected and highlighted in blue), 'student01', and 'Marketplace'. A green arrow points to the 'Marketplace' link. In the center, under the 'instructor' space, there's a table for the application 'spring-music'. The table has columns for STATUS (100%), APP (spring-music), INSTANCES (1), and MEMORY (512MB). At the top right of the central area, there's a 'Edit Space' button.

and you'll see the same service/plan/description listing in the browser:

The screenshot shows the Pivotal Web Services Marketplace page. The left sidebar is identical to the previous screenshot, with 'Marketplace' selected. The main content area is titled 'Services Marketplace' and contains a message: 'Get started with our free marketplace services. Upgrade to gain access to premium service plans.' Below this, there are eight service cards arranged in two columns of four. Each card includes the service name, icon, and a brief description. The services listed are: Searchify (Custom search you control), BlazeMeter (The JMeter Load Testing Cloud), Redis Cloud (Enterprise-Class Redis for Developers), ClearDB MySQL Database (Highly available MySQL for your Apps.), CloudAMQP (Managed HA RabbitMQ servers in the cloud), ElephantSQL (PostgreSQL as a Service), SendGrid (Email Delivery. Simplified.), and MongoLab (Fully-managed MongoDB-as-a-Service).

7.3. Creating and Binding to a Service Instance

1. Let's begin by creating a MySQL instance provided by Pivotal.

From the CLI, let's *create* a p-mysql service instance:

```
$ cf create-service p-mysql 100mb-dev spring-music-db  
Creating service spring-music-db in org CNA / space instructor as example@pivotal.io...  
OK
```

If p-mysql is not an available service, use cleardb instead (select the free spark plan: cf create-service cleardb spark spring-music-db).

2. Next we'll *bind* the newly created instance to our spring-music application:

```
$ cf bs spring-music spring-music-db  
Binding service spring-music-db to app spring-music in org ACME / space instructor as example@pivotal.io...  
OK  
TIP: Use 'cf restage' to ensure your env variable changes take effect
```

3. Notice the admonition to Use 'cf restage' to ensure your env variable changes take effect. Let's take a look at the environment variables for our application to see what's been done. We can do this by typing:

```
$ cf env spring-music
```

The subset of the output we're interested in is located near the very top, titled System-Provided:

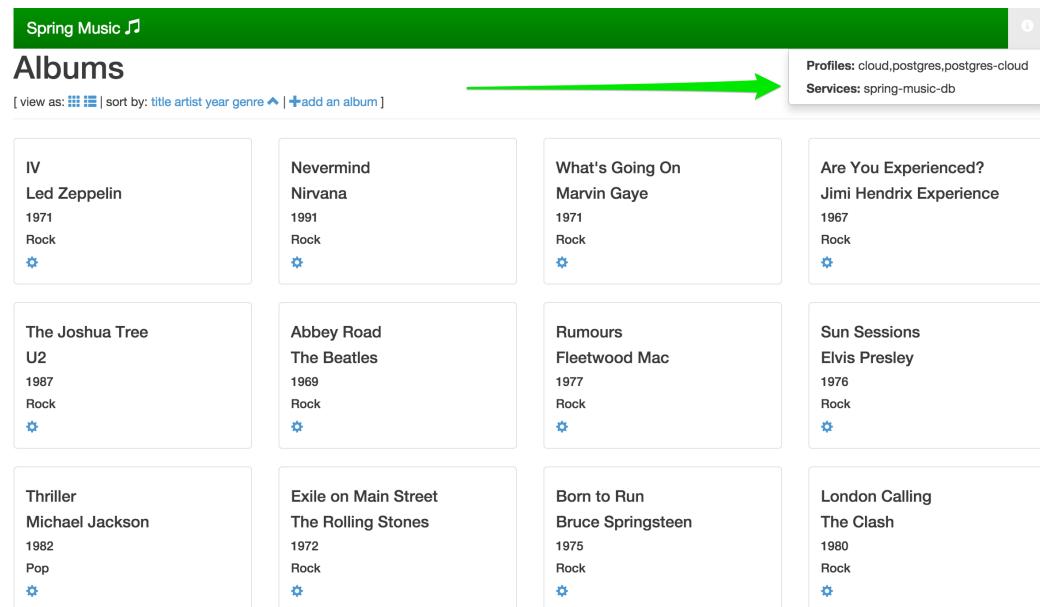
```
System-Provided:  
{  
  "VCAP_SERVICES": {  
    "p-mysql": [  
      {  
        "credentials": {  
          "hostname": "10.68.50.61",  
          "jdbcUrl": "jdbc:mysql://10.68.50.61:3306(cf_ddbe0221_dlfc_478f_b693_69279b2de702?user=mWsYAON34zbaQPAA  
\\u0026password=CZImCZ9iSNU7Hfp2",  
          "name": "cf_ddbe0221_dlfc_478f_b693_69279b2de702",  
          "password": "CZImCZ9iSNU7Hfp2",  
          "port": "3306",  
          "uri": "mysql://mWsYAON34zbaQPAA:CZImCZ9iSNU7Hfp2@10.68.50.61:3306(cf_ddbe0221_dlfc_478f_b693_69279b2de702?reconnect=true",  
          "username": "mWsYAON34zbaQPAA"  
        },  
        "label": "p-mysql",  
        "name": "spring-music-db",  
        "plan": "100mb-dev",  
        "tags": [  
          "mysql",  
          "relational"  
        ]  
      }  
    ]  
  }  
}  
  
  "VCAP_APPLICATION": {  
    "application_name": "spring-music",  
    "application_uris": [  
      "spring-music-unbragging-pourparler.pcf2.fe.gopivotal.com"  
    ],  
    "application_version": "e3fa423d-69ad-41cb-bd17-0e88427dfe4c",  
    "limits": {  
      "disk": 1024,  
      "fds": 16384,  
      "mem": 512
```

```
{
  "name": "spring-music",
  "space_id": "080bae6-bafc-482b-9a98-42f612eef736",
  "space_name": "jfullam",
  "uris": [
    "spring-music-unbragging-pourparler.pcf2.fe.gopivotal.com"
  ],
  "users": null,
  "version": "e3fa423d-69ad-41cb-bd17-0e88427dfe4c"
}
```

- ❶ VCAP_SERVICES is a special Cloud Foundry environment variable that contains a JSON document containing all of the information for any services bound to an application.
 - ❷ Notice here the unique URI for this instance of MySQL that `spring-music` has been bound to.
4. Now let's *restage* the application, which cycles our application back through the staging/buildpack process before redeploying the application.¹

```
$ cf restage spring-music
```

Once the application is running again, revisit or refresh the browser tab where you have the *Spring Music* application loaded:



The screenshot shows the Spring Music application interface. At the top, there's a navigation bar with a logo and a search bar. Below it, a header says "Albums". There are four rows of album cards. Each card displays the album title, artist, year, genre, and a settings icon. To the right of the cards, there's a sidebar with "Profiles: cloud,postgres,postgres-cloud" and "Services: spring-music-db".

Album	Artist	Year	Genre
IV	Led Zeppelin	1971	Rock
Nevermind	Nirvana	1991	Rock
What's Going On	Marvin Gaye	1971	Rock
Are You Experienced?	Jimi Hendrix Experience	1967	Rock
The Joshua Tree	U2	1987	Rock
Abbey Road	The Beatles	1969	Rock
Rumours	Fleetwood Mac	1977	Rock
Sun Sessions	Elvis Presley	1976	Rock
Thriller	Michael Jackson	1982	Pop
Exile on Main Street	The Rolling Stones	1972	Rock
Born to Run	Bruce Springsteen	1975	Rock
London Calling	The Clash	1980	Rock

¹In this case, we could accomplish the same goal by only *restarting* the application via `cf restart spring-music`. A *restage* is generally recommended because Cloud Foundry buildpacks also have access to injected environment variables and can install or configure things differently based on their values.

Click on the `(i)` symbol at top-right and it will drop down the information dialog. You should be able to see that the application is now utilizing a MySQL database via the `spring-music-db` service.

Note: If the application hasn't picked up the `spring-music-db` service, try doing a complete `cf push` again.

5. **(OPTIONAL STEPS)** If you have a MySQL GUI tool handy and you are using a lab environment that has the necessary network access (ask your instructor), you can inspect the music database created. Otherwise, your instructor will demo via a Pivotal VPN connection.
6. In your DB tool, create a new server connection and populate the properties with values from the URI in your `VCAP_SERVICES` environment variable (remember `cf env spring-music`):

7.4. Optional: Swapping from MySQL to MongoDB

Only do this next section if you have sufficient time.

1. Now let's bind our *Spring Music* application to MongoDB instead of MySQL. First let's create ² a MongoDB instance from the Pivotal Cloud Foundry marketplace:

```
$ cf cs p-mongodb development spring-music-mongo  
Creating service spring-music-mongo in org ACME / space instructor as example@pivotal.io...  
OK
```

If you do not have the `p-mongodb` service available use `mongolab` and choose the `sandbox` plan: `cf cs mongolab sandbox spring-music-mongo`.

2. Next we'll unbind our application from our MySQL instance (*Spring Music* does not support being bound to multiple datasources at the same time):

```
$ cf us spring-music spring-music-db
```

If you visit your application now, you'll see that it still works. If you recall, environment variable changes (such as binding/unbinding of services) don't actually take effect until a *restage* or *restart*.

3. Now let's bind the application to our MongoDB instance:

```
$ cf bs spring-music spring-music-mongo  
Binding service spring-music-mongo to app spring-music in org oreilly-class / space instructor as example@pivotal.io...  
OK  
TIP: Use 'cf restage' to ensure your env variable changes take effect
```

4. And then do a restage:

```
$ cf restage spring-music
```

Once the application is running again, revisit or refresh the browser tab where you have the *Spring Music* application loaded:

²Notice in this listing that we're typing `cf cs` rather than `cf create-service`. Most CF CLI commands have a shorthand version to save typing time. You can view these shorthand commands via `cf help` or `cf h` (See! More shorthand!).

Spring Music

Albums

[view as: | sort by: title artist year genre | add an album]

IV Led Zeppelin 1971 Rock 	Nevermind Nirvana 1991 Rock 	What's Going On Marvin Gaye 1971 Rock 	Are You Experienced? Jimi Hendrix Experience 1967 Rock
The Joshua Tree U2 1987 Rock 	Abbey Road The Beatles 1969 Rock 	Rumours Fleetwood Mac 1977 Rock 	Sun Sessions Elvis Presley 1976 Rock
Thriller Michael Jackson 1982 Pop 	Exile on Main Street The Rolling Stones 1972 Rock 	Born to Run Bruce Springsteen 1975 Rock 	London Calling The Clash 1980 Rock

As you can see from the information dialog, the application is now utilizing a MongoDB database via the `spring-music-mongo` service.

5. **(OPTIONAL)** Similar to the steps to view the relational data in another DB tool, a MongoDB client tool can be used to inspect the MongoDB instance. Your instructor can demo this if you do not have access to the tool and / or the appropriate network access to the MongoDB instance.

7.5. Clean Up

Because of the limited PWS quota we have for this course, let's clean up our application and services to make room for future labs.

1. Delete the `spring-music` application:

```
$ cf d spring-music
Really delete the app spring-music?> y
Deleting app spring-music in org oreilly-class / space instructor as example@pivotal.io...
OK
```

2. If you did the optional section using MongoDB, delete the `spring-music-mongo` service:

```
$ cf ds spring-music-mongo
```

```
Really delete the service spring-music-mongo?> y  
Deleting service spring-music-mongo in org oreilly-class / space instructor as example@pivotal.io...  
OK
```

3. Delete the spring-music-db service:

```
$ cf ds spring-music-db  
Really delete the service spring-music-db?> y  
Deleting service spring-music-db in org oreilly-class / space instructor as example@pivotal.io...  
OK
```

Chapter 8. Lab 2c - Scaling Applications

Cloud Foundry makes the work of horizontally scaling application instances and updating load balancer routing tables easy.

In this lab, we'll use a Spring Boot CLI app designed to illustrate Cloud Foundry operations such as scaling.

8.1. Push the cf-scale-boot Application

1. Change to the cf-scale-boot sample application directory:

```
$ cd $COURSE_HOME/lab/lab_02/lab_02c/cf-scale-boot
```

2. Spring Boot CLI applications do not require a separate build step, so go ahead and push the application:

```
$ mvn package && cf push
```

3. Once again, this application's manifest is configured to have a random route assigned to the application. So, when the CLI indicates that application is up and running, visit its route in the browser:

Cloud Foundry Scale Demo

Powered by [Spring Boot](#) and [Groovy!](#)

Kill Switch

App Instance Info:

App Name	cf-scale-boot
Instance	0
Memory Allocated	512
Disk Allocated	1024
Warden Container IP	10.254.3.86
Warden Container Port	61539
Requests Serviced	2

Figure 8.1. The application at initial deployment

You'll see that the application is reporting various bits of information that it has discovered from its environment. Of primary interest is that this application reports its *instance index* - note that this is the instance id *not* the number of instances running. It also keeps track of and reports how many web requests that this instance has serviced.

8.2. Scale the Application Up

1. Now let's increase the number of running application instances to 4 (make sure you have stopped `spring-music` before doing this or you may exceed your quota):

```
$ cf scale -i 4 cf-scale-boot
Scaling app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
OK
```

In reporting OK, the CLI is letting you know that the additional requested instances have been started, but they are not yet necessarily running.

2. We can determine how many instances are actually running like this:

```
$ cf app cf-scale-boot
Showing health and status for app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
OK

requested state: started
instances: 4/4
usage: 512M x 4 instances
urls: cf-scale-boot-stockinged-rust.cfapps.io
last uploaded: Fri Feb 13 18:56:29 UTC 2015

      state      since            cpu    memory       disk
#0  running   2015-02-13 12:57:10 PM  0.1%  404.8M of 512M  128.9M of 1G
①
#1  starting   2015-02-13 03:04:33 PM  0.0%  0 of 0        0 of 0
②
#2  running   2015-02-13 03:04:47 PM  0.0%  398.7M of 512M  128.9M of 1G
#3  starting   2015-02-13 03:04:33 PM  0.0%  0 of 0        0 of 0
```

- ❶ This application instance has completed the startup process and is actually able to accept requests.
- ❷ This application instance is still starting and will not have any requests routed to it.

3. Eventually all instances will converge to a running state:

```
$ cf app cf-scale-boot
Showing health and status for app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
OK

requested state: started
instances: 4/4
usage: 512M x 4 instances
urls: cf-scale-boot-stockinged-rust.cfapps.io
last uploaded: Fri Feb 13 18:56:29 UTC 2015

      state      since            cpu    memory       disk
#0  running   2015-02-13 12:57:10 PM  0.1%  404.8M of 512M  128.9M of 1G
```

```
#1  running  2015-02-13 03:04:51 PM  0.1%  377.5M of 512M  128.9M of 1G
#2  running  2015-02-13 03:04:47 PM  0.1%  397.3M of 512M  128.9M of 1G
#3  running  2015-02-13 03:05:03 PM  0.0%  389.2M of 512M  128.9M of 1G
```

4. Revisit the application route in the browser. Refresh several times. You should observe the instance index and request counters changing as you do so - something like this (you instance and request numbers will differ):

The screenshot shows the 'Cloud Foundry Scale Demo' application. At the top, it says 'Powered by Spring Boot and Groovy!'. Below that is a red button labeled 'Kill Switch'. The main area is titled 'App Instance Info:' and contains a table with the following data:

App Name	cf-scale-boot
Instance	3
Memory Allocated	512
Disk Allocated	1024
Warden Container IP	10.254.2.58
Warden Container Port	61142
Requests Serviced	4

Two green arrows point to the '3' under 'Instance' and the '4' under 'Requests Serviced', indicating the scaled-up state.

Figure 8.2. The application after scaling up

If scaling has worked correctly, you should see instance ids between 0 and 3.

8.2.1. About the Router

The aforementioned [\(Go\)Router](#) is applying a random routing algorithm to all of the application instances assigned to this route. As an instance reaches the `running` state, its [Execution Agent](#) (also known as a Cell) registers that instance in the routing table assigned to its route by sending a message to Cloud Foundry's message bus. All (Go)Router instances are subscribed to this channel and register the routes independently. This makes for very dynamic and rapid reconfiguration!

8.3. Scale the Application Down

1. We can scale the application instances back down as easily as we scaled them up, using the same command structure:

```
$ cf scale -i 1 cf-scale-boot  
Scaling app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...  
OK
```

2. Check the application status again:

```
$ cf app cf-scale-boot  
Showing health and status for app cf-scale-boot in org oreilly-class / space instructor  
as example@pivotal.io...  
OK  
  
requested state: started  
instances: 1/1  
usage: 512M x 1 instances  
urls: cf-scale-boot-stocked-rust.cfapps.io  
last uploaded: Fri Feb 13 18:56:29 UTC 2015  
  
#0 state      since            cpu    memory      disk  
#0  running   2015-02-13 12:57:10 PM  0.1%  405M of 512M  128.9M of 1G
```

As you can see, we're back down to only one instance running, and it is in fact the original index 0 that we started with.

3. Confirm that by again revisiting the route in the browser and checking the instance index and request counter. The instance id should now *always* be zero:

Cloud Foundry Scale Demo

Powered by [Spring Boot](#) and [Groovy!](#)

Kill Switch

App Instance Info:

App Name	cf-scale-boot
Instance	0
Memory Allocated	512
Disk Allocated	1024
Warden Container IP	10.254.3.86
Warden Container Port	61539
Requests Serviced	5

Figure 8.3. The application after scaling down

Chapter 9. Lab 2d - Monitoring Applications

Cloud Foundry provides several built-in mechanisms that allow us to monitor our applications' state changes and behavior. Additionally, Cloud Foundry actively monitors the health of our application processes and will restart them should they crash. In this lab, we'll explore a few of these mechanisms.

9.1. Events

Cloud Foundry only allows application configuration to be modified via its API. This gives application operators confidence that all changes to application configuration are known and auditable. It also reduces the number of causes that must be considered when problems arise.

All application configuration changes are recorded as *events*. These events can be viewed via the Cloud Foundry API, and viewing is facilitated via the CLI.

Take a look at the events that have transpired so far for our deployment of cf-scale-boot:

```
$ cf events cf-scale-boot
Getting events for app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
time           event          actor          description
2015-02-13T15:18:33.00-0600  audit.app.update  example@pivotal.io  instances: 1 ①
2015-02-13T15:04:34.00-0600  audit.app.update  example@pivotal.io  instances: 5 ②
2015-02-13T12:56:35.00-0600  audit.app.update  example@pivotal.io  state: STARTED ③
2015-02-13T12:56:26.00-0600  audit.app.update  example@pivotal.io  ④
2015-02-13T12:56:26.00-0600  audit.app.map-route  example@pivotal.io  ⑤
2015-02-13T12:56:24.00-0600  audit.app.create   example@pivotal.io  instances: 1, memory: 512, state: STOPPED,
environment_json: PRIVATE DATA HIDDEN ⑥
```

- ❶ Events are sorted newest to oldest, so we'll start from the bottom. Here we see the `app.create` event, which created our application's record and stored all of its metadata (e.g. `memory: 512`).
- ❷ The `app.map-route` event records the incoming request to assign a route to our application.
- ❸ This `app.update` event records the resulting change to our applications metadata.
- ❹ This `app.update` event records the change of our application's state to `STARTED`.
- ❺ Remember scaling the application up? This `app.update` event records the metadata change `instances: 5`.
- ❻ And here's the `app.update` event recording our scaling of the application back down with `instances: 1`.

1. Let's explicitly ask for the application to be stopped:

```
$ cf stop cf-scale-boot
Stopping app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
OK
```

2. Now, examine the additional `app.update` event:

```
$ cf events cf-scale-boot
Getting events for app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
time           event          actor            description
2015-02-13T15:59:10.00-0600 audit.app.update example@pivotal.io state: STOPPED
2015-02-13T15:18:33.00-0600 audit.app.update example@pivotal.io instances: 1
2015-02-13T15:04:34.00-0600 audit.app.update example@pivotal.io instances: 5
2015-02-13T12:56:35.00-0600 audit.app.update example@pivotal.io state: STARTED
2015-02-13T12:56:26.00-0600 audit.app.update example@pivotal.io
2015-02-13T12:56:26.00-0600 audit.app.map-route example@pivotal.io
2015-02-13T12:56:24.00-0600 audit.app.create  example@pivotal.io instances: 1, memory: 512, state: STOPPED,
environment_json: PRIVATE DATA HIDDEN
```

3. Start the application again:

```
$ cf start cf-scale-boot
Starting app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
0 of 1 instances running, 1 starting
0 of 1 instances running, 1 starting
0 of 1 instances running, 1 starting
1 of 1 instances running

App started

OK

App cf-scale-boot was started using this command `JAVA_HOME=$PWD/.java-buildpack/open_jdk_jre JAVA_OPTS="-Djava.io.tmpdir=$TMPDIR -XX:OnOutOfMemoryError=$PWD/.java-buildpack/open_jdk_jre/bin/killjava.sh -Xmx382293K -Xms382293K -XX:MaxMetaspaceSize=64M -XX:MetaspaceSize=64M -Xss95K" SERVER_PORT=$PORT $PWD/.java-buildpack/spring_boot_cli/bin/spring run app.groovy`

Showing health and status for app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
OK

requested state: started
instances: 1/1
usage: 512M x 1 instances
urls: cf-scale-boot-stockinged-rust.cfapps.io
last uploaded: Fri Feb 13 18:56:29 UTC 2015

      state     since          cpu    memory       disk
#0  running   2015-02-13 04:01:50 PM  0.0%  389.1M of 512M 128.9M of 1G
```

4. And again, view the additional app.update event:

```
$ cf events cf-scale-boot
Getting events for app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
time           event          actor            description
2015-02-13T16:01:28.00-0600 audit.app.update example@pivotal.io state: STARTED
2015-02-13T15:59:10.00-0600 audit.app.update example@pivotal.io state: STOPPED
2015-02-13T15:18:33.00-0600 audit.app.update example@pivotal.io instances: 1
2015-02-13T15:04:34.00-0600 audit.app.update example@pivotal.io instances: 5
2015-02-13T12:56:35.00-0600 audit.app.update example@pivotal.io state: STARTED
2015-02-13T12:56:26.00-0600 audit.app.update example@pivotal.io
2015-02-13T12:56:26.00-0600 audit.app.map-route example@pivotal.io
2015-02-13T12:56:24.00-0600 audit.app.create  example@pivotal.io instances: 1, memory: 512, state: STOPPED,
environment_json: PRIVATE DATA HIDDEN
```

9.2. Logs

One of the most important enablers of visibility into application behavior is logging. Effective management of logs has historically been very difficult. Cloud Foundry's [log aggregation](#) components simplify log management by assuming responsibility for it. Application developers need only log all messages to either STDOUT or STDERR, and the platform will capture these messages.

9.2.1. For Developers

Application developers can view application logs using the CF CLI.

1. Let's view recent log messages for cf-scale-boot:

```
$ cf logs cf-scale-boot --recent
```

Here are two interesting subsets of one output from that command:

Example 9.1. CF Component Logs

```
2015-02-13T14:45:39.40-0600 [RTR/0]      OUT cf-scale-boot-stockinged-rust.cfapps.io - [13/02/2015:20:45:39 +0000] "GET /css/bootstrap.min.css HTTP/1.1" 304 0 "http://cf-scale-boot-stockinged-rust.cfapps.io/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_5) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/40.0.2214.111 Safari/537.36" 10.10.66.88:50372 x_forwarded_for:"50.157.39.197" vcap_request_id:84cc1b7a-bb30-4355-7512-5adaf36ff767 response_time:0.013115764 app_id:7a428901-1691-4cce-b7f6-62d186c5cb55 ①
2015-02-13T14:45:39.40-0600 [RTR/1]      OUT cf-scale-boot-stockinged-rust.cfapps.io - [13/02/2015:20:45:39 +0000] "GET /img/LOGO_CloudFoundry_Large.png HTTP/1.1" 304 0 "http://cf-scale-boot-stockinged-rust.cfapps.io/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_5) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/40.0.2214.111 Safari/537.36" 10.10.66.88:2423 x_forwarded_for:"50.157.39.197" vcap_request_id:b3e2466b-6a41-4c6d-5b3d-0f70702c0ec1 response_time:0.01003444 app_id:7a428901-1691-4cce-b7f6-62d186c5cb55
2015-02-13T15:04:33.09-0600 [API/1]      OUT Tried to stop app that never received a start event ②
2015-02-13T15:04:33.51-0600 [CELL/12]    OUT Starting app instance (index 2) with guid 7a428901-1691-4cce-b7f6-62d186c5cb55 ③
2015-02-13T15:04:33.71-0600 [CELL/4]    OUT Starting app instance (index 3) with guid 7a428901-1691-4cce-b7f6-62d186c5cb55
```

- ① An “Apache-style” access log event from the (Go)Router
- ② An API log event that corresponds to an event as shown in cf events
- ③ A CELL log event indicating the start of an application instance on that execution agent (also known as a cell).

Example 9.2. Application Logs

```
2015-02-13T16:01:50.28-0600 [APP/0]      OUT 2015-02-13 22:01:50.282 INFO 36 --- [runner-0] o.s.b.a.e.jmx.EndpointMBeanExporter : Located managed bean 'autoConfigurationAuditEndpoint': registering with JMX server as MBean [org.springframework.boot:type=Endpoint,name=autoConfigurationAuditEndpoint]
2015-02-13T16:01:50.28-0600 [APP/0]      OUT 2015-02-13 22:01:50.287 INFO 36 --- [runner-0] o.s.b.a.e.jmx.EndpointMBeanExporter : Located managed bean 'shutdownEndpoint': registering with JMX server as MBean [org.springframework.boot:type=Endpoint,name=shutdownEndpoint]
2015-02-13T16:01:50.29-0600 [APP/0]      OUT 2015-02-13 22:01:50.299 INFO 36 --- [runner-0] o.s.b.a.e.jmx.EndpointMBeanExporter : Located managed bean 'configurationPropertiesReportEndpoint': registering with JMX server as MBean [org.springframework.boot:type=Endpoint,name=configurationPropertiesReportEndpoint]
2015-02-13T16:01:50.36-0600 [APP/0]      OUT 2015-02-13 22:01:50.359 INFO 36 --- [runner-0] s.b.c.e.t.TomcatEmbeddedServletContainer : Tomcat started on port(s): 61316/http
2015-02-13T16:01:50.36-0600 [APP/0]      OUT Started...
2015-02-13T16:01:50.36-0600 [APP/0]      OUT 2015-02-13 22:01:50.364 INFO 36 --- [runner-0] o.s.boot.SpringApplication : Started application in 6.906 seconds (JVM running for 15.65)
```

As you can see, Cloud Foundry’s log aggregation components capture both application logs and CF component logs relevant to your application. These events are properly interleaved based on time, giving you an accurate picture of events as they transpired across the system.

2. To get a running “tail” of the application logs rather than a dump, simply type:

```
$ cf logs cf-scale-boot
```

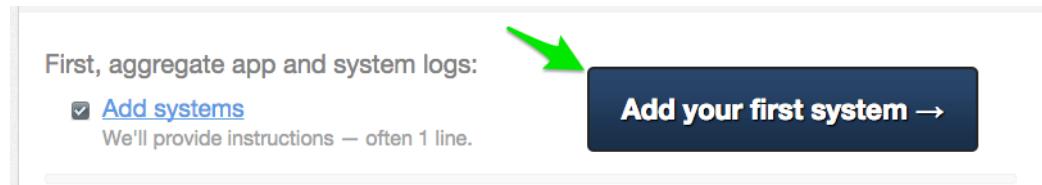
You can try various things like refreshing the browser and triggering stop/start events to see logs being generated.

9.2.2. For Operators (OPTIONAL)

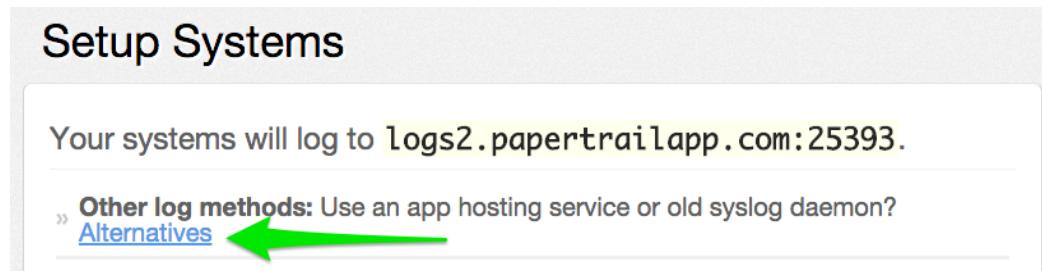
Application operators will also enjoy commands like `cf logs`, but are often interested in long-term retention, indexing, and analysis of logs as well. Cloud Foundry currently only provides short-term retention of logs. To meet these needs, Cloud Foundry provides the ability to [drain logs to third-party providers](#).

In this section, we'll drain logs to a very simple provider called [Papertrail](#).

1. Visit <https://papertrailapp.com> and create a free account.
2. Login to your account and add your first system:



3. Click on “Alternatives”:



4. Choose “I use Heroku” and provide a name:

Choose your situation:



A My syslogd only uses the default port

GNU syslogd and some embedded devices will only log to port 514. A few old Linux distro versions use GNU syslogd (mostly CentOS and Gentoo).



B I use Cloud Foundry

Register each app separately. Use Heroku? [Here's how.](#)



C My system's hostname changes

In rare cases, one system may change hostnames frequently. For example, a roaming laptop which sets its hostname based on DHCP (and roams across networks).

We'll provide an app-specific syslog drain and step-by-step setup for [Cloud Foundry](#).

Let's create a destination for this app.

What should we call it?

cf-scale-boot

Alphanumeric. Does not need to match app name.

Save 

5. Note the URL + Port assigned to your application:

Setup cf-scale-boot

 Edit Settings

System created.

cf-scale-boot will log to **logs2.papertrailapp.com:43882**.

6. We'll use a Cloud Foundry [user-provided service instance](#) to create the log drain for our application using the URL + Port provided by Papertrail:

```
$ cf cups cf-scale-boot-logs -l syslog://logs2.papertrailapp.com:43882  
Creating user provided service cf-scale-boot-logs in org oreilly-class / space instructor as example@pivotal.io...  
OK
```

7. We bind that service instance like those we created in Lab 2c:

```
$ cf bs cf-scale-boot cf-scale-boot-logs  
Binding service cf-scale-boot-logs to app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...  
OK  
TIP: Use 'cf restage' to ensure your env variable changes take effect
```

8. We'll use a `cf restart` rather than `cf restage` to make the binding take effect:

```
$ cf restart cf-scale-boot
```

9. Refresh the Papertrail "Events" tab to see log events immediately flowing to the log viewing page:

sf-scale-boot	Dashboard	Events	Account	Help	Me
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.869 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Mapped [/health],methods=[GET],params=[],headers[],consumes[],produces[],custom[] onto public java.lang.Object org.springframework.boot.actuate.endpoint.mvc.HealthMvcEndpoint.invoke()					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.874 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Mapped [/metrics],methods=[GET],params=[],headers[],consumes[],produces[],custom[] onto public java.lang.Object org.springframework.boot.actuate.endpoint.mvc.HealthMvcEndpoint.invoke()					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.870 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Mapped [/beans],methods=[GET],params=[],headers[],consumes[],produces[],custom[] onto public java.lang.Object org.springframework.boot.actuate.endpoint.mvc.HealthMvcEndpoint.invoke()					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.874 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Mapped [/dump],methods=[GET],params=[],headers[],consumes[],produces[],custom[] onto public java.lang.Object org.springframework.boot.actuate.endpoint.mvc.EndpointMvcAdapter.invoke()					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.878 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Mapped [/metrics/{name}],methods=[GET],params=[],headers[],consumes[],produces[],custom[] onto public java.lang.Object org.springframework.boot.actuate.endpoint.mvc.HealthMvcEndpoint.invoke()					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.872 INFO 36 --- [runner-0] o.s.c.support.DefaultLifecycleProcessor : Starting beans in phase 0					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.903 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [requestEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=requestMappingEndpoint]					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.906 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [tracingEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=traceMappingEndpoint]					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.946 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [healthEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=healthEndpoint]					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.951 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [beansEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=beansEndpoint]					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.968 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [informationEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=informationEndpoint]					
Feb 13 14:57:10 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.964 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [metricsEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=metricsEndpoint]					
Feb 13 14:57:11 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:11.064 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [tracingEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=traceMappingEndpoint]					
Feb 13 14:57:11 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:11.068 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [dumpEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=dumpMappingEndpoint]					
Feb 13 14:57:11 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.976 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [autoConfigurationAuditEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=autoConfigurationAuditEndpoint]					
Feb 13 14:57:11 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.983 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [shardEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=shardMappingEndpoint]					
Feb 13 14:57:11 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:10.993 INFO 36 --- [runner-0] o.s.b.a.mvc.EndpointHandlerMapping : Located managed bean [configurationPropertiesReportEndpoint]: registering with JMx server as Mbean [org.springframework.boot:type=Endpoint,name=configurationPropertiesReportEndpoint]					
Feb 13 14:57:11 cf-scale-boot 7d428901-1691-4cc6-b7f6-62d186c5b55 [/App:0] 2015-02-13 22:57:11.053 INFO 36 --- [port=5: 61617/http] [runner-0] s.b.c.e.TomcatEmbeddedServletContainer : Tomcat started on 6.882 seconds (CM running for 12.589s)					
Example: "access denied" 1.2.3.4 -sshd	Search	Contrast	Pause		

You can see how to connect to other third-party log management systems in the [Cloud Foundry documentation](#).

9.3. Health

Cloud Foundry's [Health Manager](#) actively monitors the health of our application processes and will restart them should they crash.

1. If you don't have one already running, start a log tail for `cf-scale-boot`:

```
$ cf logs cf=scale-boot
```

2. Visit the application in the browser, and click on the “Kill Switch” button. This button will trigger a JVM exit with an error code (`System.exit(1)`), causing the Health Manager to observe an application instance crash;

Cloud Foundry Scale Demo

Powered by Spring Boot and Groovy!

Kill Switch

Go ahead...CLICK IT!

App Instance Info:

App Name	cf-scale-boot
Instance	0
Memory Allocated	512
Disk Allocated	1024
Warden Container IP	10.254.1.38
Warden Container Port	61617
Requests Serviced	1

3. After clicking the kill switch a couple of interesting things should happen. First, you'll see an error code returned in the browser, as the request you submitted never returns a response:

502 Bad Gateway: Registered endpoint failed to handle the request.

Also, if you're paying attention to the log tail, you'll see some interesting log messages fly by:

```
2015-02-13T17:17:54.86-0600 [App/0]      OUT 2015-02-13 23:17:54.860 ERROR 36 --- [io-61617-exec-5] WebApplication
: KILL SWITCH ACTIVATED! ❶
2015-02-13T17:17:54.86-0600 [App/0]      OUT 2015-02-13 23:17:54.869 INFO 36 --- [           Thread-2] a.t.a.c.e.EmbeddedWeb
ebApplicationContext : Closing org.springframework.boot.c$next.embedded.AnnotationConfigEmbeddedWebApplicationContext@6a62811d:
startUp date: [Fri Feb 13 22:57:05 UTC 2015]; root of context hierarchy
2015-02-13T17:17:54.87-0600 [App/0]      OUT 2015-02-13 23:17:54.870 INFO 36 --- [           Thread-2] o.s.c.support.Defaul
tLifecycleProcessor : Stopping beans in phase 0
2015-02-13T17:17:54.87-0600 [App/0]      OUT 2015-02-13 23:17:54.874 INFO 36 --- [           Thread-2] o.s.b.a.e.jmx.Endpoi
ntMBeanExporter : Unregistering JMX-exposed beans on shutdown
2015-02-13T17:17:54.87-0600 [App/0]      OUT 2015-02-13 23:17:54.878 INFO 36 --- [           Thread-2] o.s.j.e.a.Annotation
MBeanExporter : Unregistering JMX-exposed beans on shutdown
2015-02-13T17:17:57.30-0600 [RTR/1]      OUT cf-scale-boot-stockinged-rust.cfapps.io - [13/02/2015:23:17:54 +0000] "GET
/killSwitch HTTP/1.1" 502 0 "http://cf-scale-boot-stockinged-rust.cfapps.io/" "Mozilla/5.0 (Macintosh; Intel Mac OS X
10_9_5) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/40.0.2214.111 Safari/537.36" 10.10.2.122:25194
x_forwarded_for:"50.157.39.197" vcap_request_id:"fc2b93a9-451d-460f-726e-14ada0069ff4" response_time:2.465784807
app_id:7a428901-1691-4cce-b7f6-62d186c5cb55 ❷
2015-02-13T17:17:57.31-0600 [App/0]      ERR
2015-02-13T17:17:57.38-0600 [API/2]      OUT App instance exited with guid 7a428901-1691-4cce-b7f6-62d186c5cb55 payload:
{"cc_partition"=>"default", "droplet"=>"7a428901-1691-4cce-b7f6-62d186c5cb55", "version"=>"ebcd2b62-2851-4716-83a4
-c816fa2c68bb", "instance"=>"leecfb8d3b41492a8e36237b365a4755", "index"=>0, "reason"=>"CRASHED", "exit_status"=>1,
"exit_description"=>"app instance exited", "crash_timestamp"=>1423869477} ❸
```

- ① Just before issuing the `System.exit(1)` call, the application logs that the kill switch was clicked.
 - ② The (Go)Router logs the 502 error.
 - ③ The API logs that an application instance exited due to a crash.

4. Check the application events to see another indicator of the crash:

```
$ cf events cf-scale-boot
Getting events for app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...
time           event          actor        description
2015-02-13T17:17:57.00-0600  app.crash    cf-scale-boot  index: 0, reason: CRASHED, exit_description: app
   instance exited, exit status: 1
```

5. By this time you should have noticed some additional interesting events in the logs:

- ① The CELL indicates that it is starting another instance of the application as a result of the health-management system observing a difference between the desired and actual state (i.e. running instances = 1 vs. running instances = 0).
 - ② The new application instance starts logging events as it starts up.

6. Revisiting the **HOME PAGE** of the application (don't simply refresh the browser as you're still on the `/killSwitch` endpoint and you'll just kill the application again!) and you should see a fresh instance started:

Cloud Foundry Scale Demo

Powered by [Spring Boot](#) and [Groovy!](#)

Kill Switch

App Instance Info:

App Name	cf-scale-boot
Instance	0
Memory Allocated	512
Disk Allocated	1024
Warden Container IP	10.254.3.118
Warden Container Port	61222
Requests Serviced	2

9.4. Clean Up

Because of the limited PWS quota we have for this course, let's clean up our application and services to make room for future labs.

1. Delete the `cf-scale-boot` application:

```
$ cf d cf-scale-boot  
Really delete the app cf-scale-boot?> y  
Deleting app cf-scale-boot in org oreilly-class / space instructor as example@pivotal.io...  
'OK'
```

2. Delete the `cf-scale-boot-logs` service:

```
$ cf ds cf-scale-boot-logs  
Really delete the service cf-scale-boot-logs?> y  
Deleting service cf-scale-boot-logs in org oreilly-class / space instructor as example@pivotal.io...  
OK
```

Chapter 10. Lab 03a - Build a Hypermedia-Driven RESTful Web Service with Spring Data REST

Estimated time to complete: 30 minutes

In this lab we'll utilize Spring Boot, Spring Data, and Spring Data REST to create a fully-functional hypermedia-driven RESTful web service. We'll then deploy it to Pivotal Cloud Foundry.

10.1. Initializing the Application

1. Import the project \$COURSE_HOME/lab_03/lab_03a/initial/companies into your IDE.
2. Open its pom.xml file and add a runtime dependency on the H2 in-memory database:

```
<dependency>
  <groupId>com.h2database</groupId>
  <artifactId>h2</artifactId>
  <scope>runtime</scope>
</dependency>
```

3. Create the package io.pivotal.demo.domain and in that package create the class Company. Into that file you can paste the following source code, which represents companies based on exchange and trading symbol:



Note

JPA (Java Persistence API) is an abstraction layer on top of Hibernate. When importing annotations such as @Entity, you should use the package javax.persistence instead of org.hibernate.

```
@Entity
@Table(name = "COMPANIES")
public class Company implements Serializable {
    /**
     *
     */
    private static final long serialVersionUID = 1L;

    @Id
    @GeneratedValue(strategy = GenerationType.AUTO)
    private Long id;

    private String name;

    private String symbol;

    private String exchange;

    public Long getId() {
        return id;
    }

    public void setId(Long id) {
        this.id = id;
    }
}
```

Lab 03a - Build a Hypermedia-Driven RESTful Web Service with Spring Data REST

```
}

public String getName() {
    return name;
}

public void setName(String name) {
    this.name = name;
}

public String getSymbol() {
    return symbol;
}

public void setSymbol(String symbol) {
    this.symbol = symbol;
}

public String getExchange() {
    return exchange;
}

public void setExchange(String exchange) {
    this.exchange = exchange;
}
}
```

1. Create the package `io.pivotal.demo.repositories` and in that package create the interface `CompanyRepository`. Paste the following code and add appropriate imports:

```
@RepositoryRestResource(collectionResourceRel = "companies", path = "companies")
public interface CompanyRepository extends PagingAndSortingRepository<Company, Long> {
```

2. Add JPA and REST Repository support to the `io.pivotal.demo.quotes.CompaniesApplication` class that was generated by Spring Initializr.

```
@SpringBootApplication
@EnableJpaRepositories // <---- Add this
@Import(RepositoryRestMvcConfiguration.class) // <---- And this
public class CompaniesApplication {

    public static void main(String[] args) {
        SpringApplication.run(QuotesApplication.class, args);
    }
}
```

3. Run the application and open <http://localhost:8080/companies> in your browser.

You'll see that the primary endpoint automatically exposes the ability to page, size, and sort the response JSON.

So what have you done? Created four small classes and one build file, resulting in a fully-functional REST microservice. The application's `DataSource` is created automatically by Spring Boot using the in-memory database because no other `DataSource` was detected in the project.

```
{
    "_embedded" : {
        "companies" : [ ]
    },
    "_links" : {
        "self" : {
            "href" : "http://localhost:8080/companies"
        },
        "profile" : {
            "href" : "http://localhost:8080/profile/companies"
        }
    },
    "page" : {
        "size" : 20,
        "totalElements" : 0,
        "totalPages" : 0,
    }
}
```

```
{
    "number" : 0
}
```

Next we'll import some data.

10.2. Importing Data

1. Inside your `'initial'` folder, there is a file called `import.sql`. Add it to `src/main/resources`. This is a rather large dataset containing all of the companies and their symbol codes in the NASDAQ and NYSE exchanges. This file will automatically be picked up by Hibernate and imported into the in-memory database.



Note

Here we simply use a Spring Boot convention that a file named `import.sql` in the root of the classpath will be executed at startup time. You can read [here](#) if you'd like to know more details about it.

1. Run the application and access it again in your browser
2. Notice the appropriate hypermedia is included for `next`, `previous`, and `self`. You can also select pages and page size by utilizing `?size=n&page=n` on the URL string. Finally, you can sort the data utilizing `?sort=fieldName`.

```
{
    "_embedded" : {
        "companies" : [ {
            "name" : "1347 Capital Corp.",
            "symbol" : "TFSC",
            "exchange" : "NASDAQ",
            "_links" : {
                "self" : {
                    "href" : "http://localhost:8080/companies/1"
                },
                "company" : {
                    "href" : "http://localhost:8080/companies/1"
                }
            }
        },
        ...
    },
    {
        "name" : "A V Homes, Inc.",
        "symbol" : "AVHI",
        "exchange" : "NASDAQ",
        "_links" : {
            "self" : {
                "href" : "http://localhost:8080/companies/20"
            },
            "company" : {
                "href" : "http://localhost:8080/companies/20"
            }
        }
    }
],
"_links" : {
    "first" : {
        "href" : "http://localhost:8080/companies?page=0&size=20"
    }
},
```

```
    "self" : {
      "href" : "http://localhost:8080/companies"
    },
    "next" : {
      "href" : "http://localhost:8080/companies?page=1&size=20"
    },
    "last" : {
      "href" : "http://localhost:8080/companies?page=320&size=20"
    },
    "profile" : {
      "href" : "http://localhost:8080/profile/companies"
    }
  },
  "page" : {
    "size" : 20,
    "totalElements" : 6419,
    "totalPages" : 321,
    "number" : 0
  }
}
```

3. Try to access the following urls to see how the application behaves:

- <http://localhost:8080/companies?size=5>
- <http://localhost:8080/companies?size=5&page=3>
- <http://localhost:8080/companies?sort=symbol.desc>

Next we'll add searching capabilities.

10.3. Adding Search

1. Let's add some additional finder methods to CompanyRepository:

```
@RestResource(path = "name", rel = "name")
Page<Company> findByNameIgnoreCase(@Param("q") String name, Pageable pageable);

@RestResource(path = "nameContains", rel = "nameContains")
Page<Company> findByNameContainsIgnoreCase(@Param("q") String name, Pageable pageable);

@RestResource(path = "symbol", rel = "symbol")
Page<Company> findBySymbolIgnoreCase(@Param("q") String symbol, Pageable pageable);

@RestResource(path = "exchange", rel = "exchange")
Page<Company> findByExchange(@Param("q") String exchange, Pageable pageable);
```

2. Run the application

3. Access the application again. Notice that hypermedia for a new search endpoint has appeared at the end of the page.

```
{
  ...
  "_links" : {
    "next" : {
      "href" : "http://localhost:8080/companies?page=1&size=20"
    },
    "self" : {
      "href" : "http://localhost:8080/companies{?page,size,sort}",
      "templated" : true
    },
    "search" : {
      "href" : "http://localhost:8080/companies/search"
    }
  }
}
```

```
},  
...  
}
```



Note

The search endpoint is automatically added by Spring Data REST when you have declared at least one `findBy` method

1. Access the new search endpoint on <http://localhost:8080/companies/search>

```
{  
  "_links" : {  
    "nameContains" : {  
      "href" : "http://localhost:8080/companies/search/nameContains{?q,page,size,sort}",  
      "templated" : true  
    },  
    "symbol" : {  
      "href" : "http://localhost:8080/companies/search/symbol{?q,page,size,sort}",  
      "templated" : true  
    },  
    "name" : {  
      "href" : "http://localhost:8080/companies/search/name{?q,page,size,sort}",  
      "templated" : true  
    },  
    "exchange" : {  
      "href" : "http://localhost:8080/companies/search/exchange{?q,page,size,sort}",  
      "templated" : true  
    },  
    "Self" : {  
      "href" : "http://localhost:8080/companies/search"  
    }  
  }  
}
```

Note that we now have new search endpoints for each of the finders that we added.

2. Try a few of these endpoints. Feel free to substitute your own values for the parameters.

<http://localhost:8080/companies/search/exchange?q=NASDAQ>

<http://localhost:8080/companies/search/name?q=Amerco>

<http://localhost:8080/companies/search/nameContains?q=Apple&size=5>

10.4. Pushing to Cloud Foundry

1. An empty `manifest.yml` file has been placed at the root of your project. Paste the below configuration inside your manifest file:

```
---  
applications:  
- name: companies  
  host: companies-${random-word}  
  memory: 512M  
  instances: 1  
  path: target/lab_03a-companies-0.0.1-SNAPSHOT.jar  
  timeout: 180 # to give time for the data to import
```

2. Build the JAR and push it to Cloud Foundry (you need to leave the IDE to do this):

Lab 03a - Build a Hypermedia-Driven RESTful Web Service with Spring Data REST

```
mvn package  
cf push
```

3. Once it has uploaded and run the buildpack, you should see output like this ...

```
1 of 1 instances running  
  
App started  
  
OK  
  
App companies was started using this command `CALCULATED_MEMORY=$(($PWD/.java-buildpack/open_jdk_jre/bin/java-buildpack  
-memory-calculator-2.0.0_RELEASE -memorySizes=metaspace:64m.. -memoryWeights=heap:75,metaspace:10,native:10,stack:5  
-memoryInitials=heap:100%,metaspace:100% -totMemory=$MEMORY_LIMIT) && SERVER_PORT=$PORT $PWD/.java-buildpack  
/open_jdk_jre/bin/java -cp $PWD/.java-buildpack/spring_auto_reconfiguration/spring_auto_reconfiguration  
-1.10.0_RELEASE.jar -Djava.io.tmpdir=$TMPDIR -XX:OnOutOfMemoryError=$PWD/.java-buildpack/open_jdk_jre/bin/killjava.sh  
$CALCULATED_MEMORY org.springframework.boot.loader.JarLauncher'  
  
Showing health and status for app companies in org pivot-cqueiroz / space development as cqueiroz@pivotal.io...  
OK  
  
requested state: started  
instances: 1/1  
usage: 512M x 1 instances  
uri: companies-getable-section.cfapps.pivotal.io  
last uploaded: Thu Nov 26 11:02:33 UTC 2015  
stack: cflinuxfs2  
buildpack: java-buildpack=v3.3.1-offline-https://github.com/cloudfoundry/java-buildpack.git#063836b java-main open-jdk  
-like-jre=1.8.0_65 open-jdk-like-memory-calculator=2.0.0_RELEASE spring-auto-reconfiguration=1.10.0_RELEASE  
  
#0 state      since            cpu      memory         disk       details  
#0  running   2015-11-26 07:03:09 PM  0.0%   442.3M of 512M  151.1M of 1G
```

4. Access the application at the random route provided by CF - your URL will be different, look in the App Manager console to see what it is:

```
http://companies-<your_path>.cfapps.io/companies
```

Chapter 11. Lab 03b - Leveraging Spring Cloud Connectors for Service Binding

In this lab we'll bind our RESTful web service from the [previous lab](#) to a MySQL database and leverage Spring Cloud Connectors to easily connect to it.



Note

The completed code for this lab can be found at
\$COURSE_HOME/lab_03/lab_03b/solution/companies.

11.1. Using Spring Cloud Connectors

1. Change to the lab directory (the initial state for this lab is the same as the completed state for the [previous lab](#), so you can choose to continue with that project if you like):

```
$ cd $COURSE_HOME/lab/lab_03/lab_03b/initial/companies
```

2. At present we're still using the in-memory database. Let's connect to a MySQL database service. From the CLI, let's *create* a MySQL service instance:

```
$ cf cs p-mysql 100mb-dev companies-db  
Creating service companies-db...  
OK
```

Again, if p-mysql service is not available, please use the free spark plan from cleardb like this cf cs cleardb spark companies-db.

3. Next add the service to your application manifest, which will *bind* the service to our application on the next push. We'll also add an environment variable to switch on the "cloud" profile,

```
---  
applications:  
- name: companies  
  memory: 512M  
  instances: 1  
  path: target/lab_03b-companies-1.0.0.jar  
  timeout: 180  
  services:  
    - companies-db # Add  
    # these  
  env:  
    SPRING_PROFILES_ACTIVE: cloud # Four lines
```

You can also accomplish the service binding by explicitly binding the service at the command-line:

```
$ cf bind-service companies companies-db  
Binding service companies-db to app companies...  
OK
```

4. Next we'll add Spring Cloud and MySQL dependencies to our maven pom file. Just after the dependency

declaration for h2, add add the following in the dependencies section:

```
<dependency>
    <groupId>org.springframework.cloud</groupId>
    <artifactId>spring-cloud-spring-service-connector</artifactId>
</dependency>
<dependency>
    <groupId>org.springframework.cloud</groupId>
    <artifactId>spring-cloud-cloudfoundry-connector</artifactId>
</dependency>
```

5. Next, let's create the package `io.pivotal.demo.config` and create in that package the class `CloudDataSourceConfig`. Add the following code:

```
@Profile("cloud")
@Configuration
public class CloudDataSourceConfig extends AbstractCloudConfig {
    @Bean
    public DataSource dataSource() {
        return connectionFactory().dataSource();
    }
}
```

The `@Profile` annotation will cause this class (which becomes Spring configuration when annotated as `@Configuration`) to be added to the configuration set because of the `SPRING_PROFILES_ACTIVE` environment variable we added earlier. You can still run the application locally (with the default profile) using the embedded database.

With this code, Spring Cloud will detect a bound service that is compatible with `DataSource`, read the credentials, and then create a `DataSource` as appropriate (it will throw an exception otherwise).

6. Add the following to `src/main/resources/application.yml` to cause Hibernate to create the database schema and import data at startup. This is done automatically for embedded databases, but not for a custom `DataSource`. Other Hibernate native properties can be set in a similar fashion:

```
spring.jpa:
  hibernate.ddl-auto: create
```

7. Build the application:

```
$ mvn clean package
```

8. Re-push the application:

```
$ cf push
```

9. In your web browser, have a look at the env variables used by your application (http://companies-<your_path>.cfapps.io/env)

```
...
"vcap": {
  "vcap.services.companies-db.name": "companies-db",
  "vcap.application.limits.mem": "1024",
  "vcap.services.companies-db.label": "p-mysql",
  ...
  "vcap.services.companies-db.tags": "mysql,relational",
  ...
```

The application is now running against a MySQL database.

11.2. Customising the Datasource

1. You can customize the database connection that Spring Cloud creates with a few lines of code. Change the `dataSource` method in `CloudDataSourceConfig` to add some pooling and connection configuration:

```
@Bean  
public DataSource dataSource() {  
    PooledServiceConnectorConfig.PoolConfig poolConfig =  
        new PooledServiceConnectorConfig.PoolConfig(5, 200);  
  
    DataSourceConfig.ConnectionConfig connectionConfig =  
        new DataSourceConfig.ConnectionConfig("characterEncoding=UTF-8");  
    DataSourceConfig serviceConfig = new DataSourceConfig(poolConfig, connectionConfig);  
  
    return connectionFactory().dataSource("companies-db", serviceConfig);  
}
```

2. Build the application:

```
$ mvn clean package
```

3. Re-push the application:

```
$ cf push
```

Chapter 12. Lab 04a - Identify the bounded contexts of the Monolith App

In this lab we'll identify the bounded contexts so that they can be easily extracted into microservices.

12.1. Understanding the spring-trader monolith application

1. Change to the lab directory:

```
$ cd $COURSE_HOME/lab/lab_04/lab_04a/springtrader-monolith
```

and import the project (via `pom.xml`) into your IDE of choice.

2. Run the application

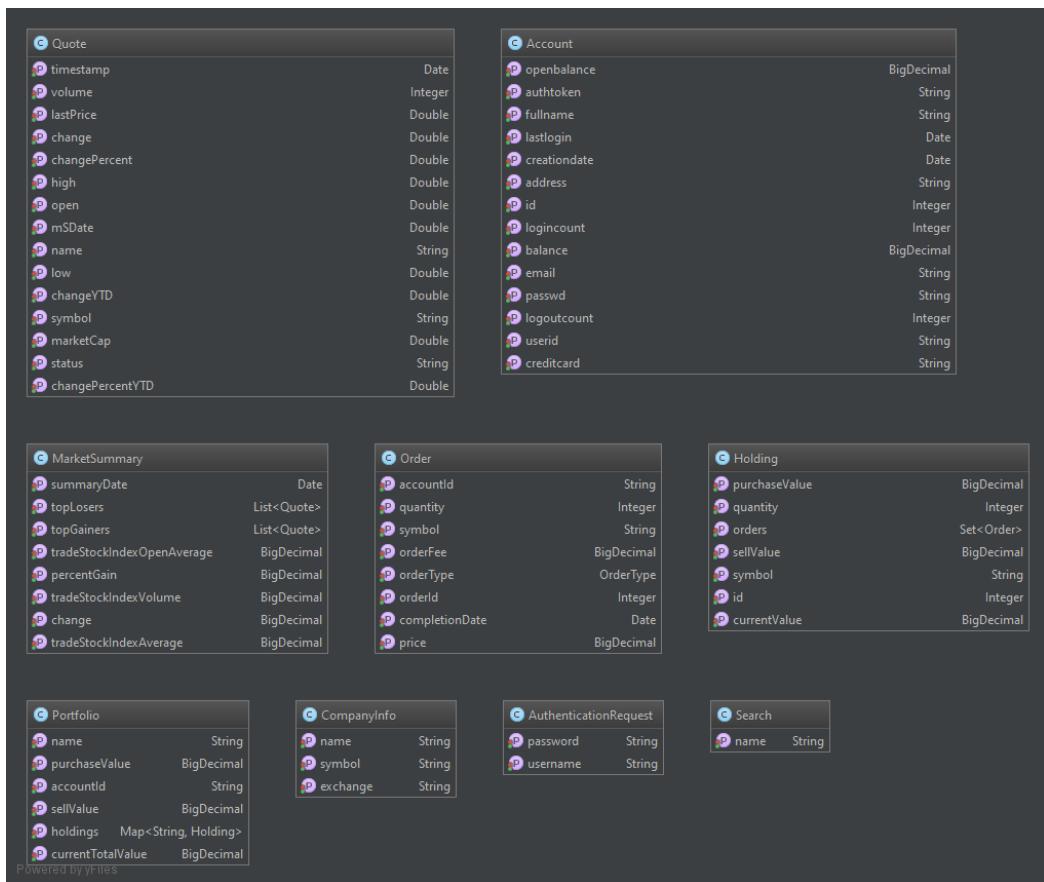
3. Browse spring-trader on <http://localhost:8080>

Create a user account with at least \$5000 balance, perform a few trades and check your portfolio.

4. Domain classes

Inside your IDE, look at the domain classes (package `io.pivotal.springtrader.domain`).

Try to write down all the relations between those domain classes



12.2. Decomposition into microservices

1. Based on the domain model, identify how the Spring Trader monolith application should be split up.

Chapter 13. Lab 05a - Build a Quotes services using MongoDB

Estimated time to complete: 30 minutes

In this lab we'll begin the first subproject of our complete cloud-native application architecture: *SpringTrader*, which implements the quotes fetching portion of a '*Trading-like*' application.

This microservice will provide the *Quotes Service*, which will allow us to search for companies and symbols in order to make transactions.

We'll be using the [MarkitOnDemand](#) API, and we'll model its `Quote` and `Company` as a `Stock` concept in this microservice.

The application that we are building is structured like this:

```
      +---- Quotes
      |
Web UI  ----- Accounts
      |
      +--- Portfolio
```

The next four labs look at each component in turn, starting with Quotes.

13.1. Exploring springtrader-quotes

1. Change to the lab directory:

```
$ cd $COURSE_HOME/lab/lab_05/lab_05a/initial/springtrader-quotes
```

and import the project (via `pom.xml`) into your IDE.

2. Explore the new `quotes` microservice:

- How many domain classes do you now have?
- How many controller classes?
- You can see that `html`, `css` and `javascript` files are not included in `springtrader-quotes`. They will all be contained in the `WebUI` microservice (a dedicated web front-end application) that we will see in a later lab.
- This is a REST application returning market data in JSON format.

- Stock data is stored in MongoDB via the `StockRepository` - if you examine the code you can see it extends `MongoRepository` and uses Spring Data to auto-generate the code
3. Run the application from your IDE (if using STS run as a Spring Boot application) It defaults to using an embedded version of MongoDB
4. Access the application at <http://localhost:8086/quote/msft>
You should see output like this:

```
{  
  "Symbol": "MSFT",  
  "Name": "Microsoft Corp",  
  "Exchange": "NASDAQ",  
  "Status": "SUCCESS",  
  "LastPrice": 53.95,  
  "Change": 0.260000000000005,  
  "ChangePercent": 0.484261501210663,  
  "Timestamp": "Fri Nov 27 17:59:00 GMTZ 2015",  
  "MSDate": 42335.5409722222,  
  "MarketCap": 4.3094790635E11,  
  "Volume": 716782,  
  "ChangeYTD": 46.45,  
  "ChangePercentYTD": 16.1463939720129,  
  "High": 54.08,  
  "Low": 53.81,  
  "Open": 53.94  
}
```

13.2. Preparing for Cloud Foundry

- We want to bind to a MongoDB data source when running on Cloud Foundry, so let's create one:
We will use the `cf create-service` command, which can be abbreviated to `cf cs`:

```
$ cf cs p-mongodb development springtrader-quotes-db  
Creating service springtrader-quotes-db in org cqueiroz-pivot / space development as cqueiroz@pivotal.io...  
OK
```



Note

If you are using Pivotal Web Service, the above service is not available in the command line. You just need to create it using the Web console at <https://console.run.pivotal.io>. You can select a MongoDB service and name it `springtrader-quotes-db`.

- Go into the package `io.pivotal.springtrader.quotes.config` and observe the differences between `MongoCloudConfig` (cloud profile) and `MongoLocalConfig` (local profile)
As you will see, the local profile provides you with a MongoDB in-memory database. The cloud profile will rely on the service that you have just created on Cloud Foundry.
- Open `src/main/resources/application.yml` and see that a `cloud` profile has been setup in preparation for

deployment to Cloud Foundry:

```
...  
spring:  
  profiles: cloud  
...  
...
```

3. Rebuild the JAR:

```
$ mvn clean package
```

13.3. Deploying to Cloud Foundry

1. Open your `manifest.yml` file and paste inside the following configuration:

```
timeout: 180  
instances: 1  
memory: 512M  
env:  
  SPRING_PROFILES_ACTIVE: cloud  
  JAVA_OPTS: -Djava.security.egd=file:///dev/urandom  
applications:  
- name: quotes  
  random-route: true  
  path: target/lab_05a-quotes-1.0.0-SNAPSHOT.jar  
  services: [ springtrader-quotes-db ]
```

2. Push to Cloud Foundry. Assuming it pushes successfully, the health and status output will show you the URL your application has been mapped to (`urls` in the output below):

```
$ cf push  
...  
Showing health and status for app quotes in org pivot-cqueiroz / space development as cqueiroz@pivotal.io...  
OK  
requested state: started  
instances: 1/1  
usage: 512M x 1 instances  
urls: quotes-undespairing-1enision.cfapps.petz.pivotai.io  
last uploaded: Mon Dec 7 22:17:58 UTC 2015  
stack: clinuxfs2  
buildpack: java-buildpack=v3.3.1-offline-https://github.com/cloudfoundry/java-buildpack.git#063836b java-main java-opts  
  open-jdk-like-jre=1.8.0_65 open-jdk-like-memory-calculator=2.0.0_RELEASE spring-auto-reconfiguration=1.10.0_RELEASE  
state      since            cpu      memory      disk      details  
#0  running   2015-12-07 07:18:48 PM  3.7%  367.2M of 512M  140.2M of 1G
```

3. Using `curl` or your browser, access the application at the random route provided by CF (in this case `http://quotes-xxxx-yyyy.cfapps.io/quote/aapl`). The output should look something like:

```
{  
  "Symbol": "AAPL",  
  "Name": "Apple Inc.",  
  "Exchange": "NASDAQ",  
  "Status": "SUCCESS",  
  "LastPrice": 117.82,  
  "Change": -0.2100000000000008,  
  "ChangePercent": -0.17792086756047,  
  "Timestamp": "Fri Nov 27 17:59:00 GMT+02015",  
  "MSDate": "42335.5409722222,  
  "MarketCap": 6.5688549842E11,  
  "Volume": 859197,
```

```
"ChangeYTD":110.38,  
"ChangePercentYTD":6.74035151295524,  
"High":118.41,  
"Low":117.6,  
"Open":118.27  
}
```

Chapter 14. Lab 05b - Build Accounts Service with Mysql/H2

Estimated time to complete: 30 minutes

In this lab we'll begin the second subproject of our complete cloud-native application architecture: *SpringTrader*.

This microservice will provide the *Accounts Service*, which will allow us to create and browse user accounts. We'll be using a relational model to represent the user accounts.



Note

The completed code for this lab can be found at

`$COURSE_HOME/lab_05/lab_05b/solution/springtrader-accounts.`

14.1. Exploring springtrader-accounts

1. Change to the lab directory:

```
$ cd $COURSE_HOME/lab/lab_05/lab_05b/initial/springtrader-accounts
```

and import the project (via `pom.xml`) into your IDE of choice.

2. Explore the new `accounts` microservice:

- How many domain classes do you now have? Are there relationships between those classes?

- How many controller classes?

Like `spring-trader-quotes`, you can see that static web files are not included in `springtrader-accounts`. This is also a REST application returning JSON data

3. Review the `application.yml`, it is empty. What do we need?

- To run multiple apps locally we will need to define a port other than 8080
- To use JPA to persist account data into our RDBMS, we need to help Spring Boot configure it
- To set the logging level to INFO - if you have problems later, you can change INFO to DEBUG
Copy this code ...

```
server.port: 8082
security.basic.enabled: false
```

```
spring.jpa:  
  hibernate.ddl-auto: create  
  
logging:  
  level:  
    io.pivotal.springtrader: 'INFO'
```

1. To run the application locally we will use an in-memory H2 database which is sufficient for testing. You need to add the H2 dependency to your pom.xml file.

```
<dependency>  
  <groupId>com.h2database</groupId>  
  <artifactId>h2</artifactId>  
  <scope>runtime</scope>  
</dependency>
```

2. Open src/main/resources/import.sql. As seen previously, Spring Boot automatically loads this file at startup time (based on a naming convention). You will see there are two accounts which will be created by default at application startup time.
3. Run the application by launching AccountsApplication and open the url <http://localhost:8082/account/1> into your web browser. You should see the following:

```
{  
  "id":1,  
  "address":"45 Test Dr.",  
  "passwd":"123",  
  "userid":"johnsmith",  
  "email": "john@pivotal.io",  
  "creditcard": "9999999999",  
  "fullname": "John Smith",  
  "authToken": null,  
  "creationDate": null,  
  "openBalance": 1200.50,  
  "logoutCount": 10,  
  "balance": null,  
  "lastLogin": null,  
  "loginCount": 0  
}
```

14.2. Preparing for Cloud Foundry

1. Create the MySQL service.

We will use the cf create-service command, which can be abbreviated to cf cs:

```
$ cf cs p-mysql 100mb-dev springtrader-accounts-db
```



Note

if you are using Pivotal Web Services (PWS), the above service is not available in the command line. You just need to create it using the Web console at <https://console.run.pivotal.io>. You can select a MySql database service and name it springtrader-accounts-db.

2. Rebuild the JAR:

```
$ mvn clean package
```

14.3. Deploying to Cloud Foundry

1. Paste the following inside `manifest.yml`:

```
---  
timeout: 180  
instances: 1  
memory: 512M  
env:  
  SPRING_PROFILES_DEFAULT: cloud  
  JAVA_OPTS: -Djava.security.egd=file:///dev/urandom  
applications:  
- name: accounts  
  random-route: true  
  path: target/lab_05b-accounts-1.0.0-SNAPSHOT.jar  
  services: [ springtrader-accounts-db ]
```

2. Push to Cloud Foundry:

```
$ cf push  
...  
Showing health and status for app accounts in org pivot-cqueiroz / space development as cqueiroz@pivotal.io...  
OK  
  
requested state: started  
instances: 1/1  
usage: 512M x 1 instances  
urls: accounts-unexploited-boneset.cfapps.io  
last uploaded: Mon Dec 7 21:01:55 UTC 2015  
stack: cflinuxfs2  
buildpack: java-buildpack=v3.3.1-offline-https://github.com/cloudfoundry/java-buildpack.git#063836b java-main java-opts  
  open-jdk-like-jre=1.8.0_65 open-jdk-like-memory-calculator=2.0.0_RELEASE spring-auto-reconfiguration=1.10.0_RELEASE  
  
      state     since            cpu      memory      disk       details  
#0   running   2015-12-07 06:02:51 PM  0.9%  456.4M of 512M  163.4M of 1G
```

3. Using curl or your browser, access the application at the random route provided by CF (in this case <http://accounts-unexploited-boneset.cfapps.io/account/1>). The output should be the same as when run locally:

```
{  
  "id":1,  
  "address":"45 Test Dr.",  
  "passwd":"123",  
  "userid":"johnsmith",  
  "email":"john@pivotal.io",  
  "creditcard":"9999999999",  
  "fullname":"John Smith",  
  "authtoken":null,  
  "creationdate":null,  
  "openbalance":1200.50,  
  "logoutcount":0,  
  "balance":null,  
  "lastlogin":null,  
  "logincount":0  
}
```

Chapter 15. Lab 05c - Build a Portfolio with MySQL

Estimated time to complete: 30 minutes

In this lab we'll begin the third subproject of our complete cloud-native application architecture: *SpringTrader*, which implements the quotes fetching portion of a “Trading-like” application. This microservice will provide the *Portfolio Service*, which will allow us to see our trading transactions and assets.



Note

The completed code for this lab can be found at

`$COURSE_HOME/lab_05/lab_05c/solution/springtrader-portfolio.`

15.1. Exploring springtrader-portfolio

1. Change to the lab directory:

```
$ cd $COURSE_HOME/lab/lab_05/lab_05c/initial/springtrader-portfolio
```

and import the project (via `pom.xml`) into your IDE of choice.

2. Explore the new `portfolio` microservice:

- How many domain classes do you now have? Look at the relationship between `Portfolio`, `Holder`, `Order`
- Open the class `PortfolioService`. See how it relies on the `quotes` and `accounts` microservices? As a consequence, the `portfolio` microservice won't work if the other two microservices are not up and running.

3. Add the following to `application.yml`. As you can see, each of our microservices runs on a different port so they can all run on your local computer.

```
spring:
  profiles.active: local
  application:
    name: portfolio

  spring.jpa.hibernate.ddl-auto: create

  logging:
    level:
      io.pivotal.springtrader: 'INFO'

  ---
  spring:
    profiles: local

  server:
    port: 8081

  quotes:
    url: http://localhost:8086
```

```
accounts:  
  url: http://localhost:8082  
  
---  
spring:  
  profiles: cloud  
  
quotes:  
  url: http://quotes-nonliteral-nautilus.cfapps.io  
  
accounts:  
  url: http://accounts-dormant-dewlap.cfapps.io
```

4. Open `src/main/resources/import.sql`. As seen before, Spring Boot automatically loads this file at startup time. You will therefore have one order created at application startup time.
5. To run the application locally you have to add H2 dependency to your `pom.xml` file.

```
<dependency>  
  <groupId>com.h2database</groupId>  
  <artifactId>h2</artifactId>  
  <scope>runtime</scope>  
</dependency>
```

6. Make sure the `accounts` and `quotes` microservices are still running
7. Run the application and open the url <http://localhost:8081/portfolio/johnsmith> in your web browser. You should see the following:

```
{  
  "accountId": "johnsmith",  
  "name": "JohnSmith",  
  "currentTotalValue": -255.45,  
  "purchaseValue": 0,  
  "sellValue": 100.00,  
  "holdings": [  
    {"EMC": {"id": null,  
            "symbol": "EMC",  
            "quantity": -10,  
            "purchaseValue": 0,  
            "sellValue": 100.00},  
    "orders": [  
      {  
        "orderId": 1,  
        "accountId": "johnsmith",  
        "symbol": "EMC",  
        "orderFee": 1.00,  
        "completionDate": 1481468400000,  
        "orderType": "SELL",  
        "price": 10.00,  
        "quantity": 10  
      }]  
    },  
    "currentValue": 25.545  
  ]  
}
```

15.2. Preparing for Cloud Foundry

1. We want to bind to a MongoDB data source when running on Cloud Foundry, so let's create one:

```
$ cf cs p-mysql 100mb-dev springtrader-portfolio-db  
Creating service instance springtrader-portfolio-db in org pivot-cqueiroz / space development as cqueiroz@pivotal.io...  
OK
```

2. Add the appropriate dependencies for the Spring Cloud Connectors:

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-cloudfoundry-connector</artifactId>
</dependency>
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-spring-service-connector</artifactId>
</dependency>
```

3. Build the JAR:

```
$ mvn clean package
```

15.3. Deploying to Cloud Foundry

1. Create an application manifest in `manifest.yml`:

```
timeout: 180
instances: 1
memory: 512M
env:
  SPRING_PROFILES_ACTIVE: cloud
  JAVA_OPTS: -Djava.security.egd=file:///dev/urandom
applications:
- name: portfolio
  random-route: true
  path: target/lab_05c-portfolio-1.0.0-SNAPSHOT.jar
  services: [ springtrader-portfolio-db ]
```

2. Before push to Cloud Foundry adjust the `application.yml` (cloud profile) file with the URLs for the Accounts and Quotes services. Something like:

```
...
...
spring:
  profiles: cloud
  jpa:
    hibernate:
      ddl-auto: create
  quotes:
    url: http://quotes-random-words1.cfapps.io
  accounts:
    url: http://accounts-random-words2.cfapps.io
```

3. Push to Cloud Foundry:

```
$ cf push
...
```

4. Access the application on <http://portfolio<your-random-words>.cfapps.io/portfolio/johnsmith>. You should see the following:

```
{
  "accountId": "johnsmith",
  "name": "JohnSmith",
  "currentTotalValue": 255.450,
  "purchaseValue": 100.00,
  "sellValue": 0,
  "holdings": {
```

```
"EMC":{  
    "id":null,  
    "symbol":"EMC",  
    "quantity":10,  
    "purchaseValue":100.00,  
    "sellValue":10,  
    "orders": [  
        {  
            "orderId":1,  
            "accountId":"johndoe",  
            "symbol": "EMC",  
            "orderFee":1.00,  
            "completionDate":1329759342904,  
            "orderType": "BUY",  
            "price":10.00,  
            "quantity":10  
        }  
    ],  
    "currentValue":25.545  
}
```

Chapter 16. Lab 05d - Run the Web User Interface

Estimated time to complete: 30 minutes

In this lab we'll finish our complete cloud-native application architecture: *SpringTrader*, by running the Web User Interface (UI). This application is already complete and will delegate to the three microservices you have worked on in the rest of this session .



Note

The code for this lab can be found at:
\$COURSE_HOME/lab_05/lab_05d/initial/spring-trader-web.

16.1. Exploring springtrader-portfolio

Please make sure all three microservices from the previous three labs - quotes, accounts and portfolios are still running. Test this by opening each of these in your browser:

1. <http://localhost:8086/quote/msft>
2. <http://localhost:8082/account/1>
3. <http://localhost:8081/portfolio/johnsmith>
4. Import the project `spring-trader-web` into your IDE of choice.
5. Run as a Java application by launching the class `WebApplication`
6. Once it is running, open your browser to <http://localhost:8080>

You should see something like this. If there are no stock quotes showing, the system is not working.

The screenshot shows the SpringTrader application interface. At the top, there are two tables: one for IT Vendors and one for Financial Services organizations. The IT Vendors table includes columns for Symbol, Price, Change, High, and Low. The FS Organizations table includes columns for Symbol, Price, Change, High, and Low. Below these tables is a login form with fields for User Name and Password, and a 'Sign In' button. To the right of the login form is a registration box with a 'Create One' link.

..! IT VEndors				
Symbol	Price	Change	High	Low
CSCO	27.33	.09 ↑	27.48	27.28
ORCL	39.47	-.13 ↓	39.8	39.16
EMC	25.545	.21 ↑	25.57	25.21

..! FS Organisations				
Symbol	Price	Change	High	Low
JPM	64.07	2.41 ↑	64.17	61.72
MS	35.27	.97 ↑	35.345	34.65
BAC	17.805	.43 ↑	17.81	17.49

Login
Enter Username and Password to login

User Name :

Password:

Sign In

Registration
Don't have a trader account yet?
[Create One](#)

16.2. Running the Application

1. The first thing you need to do is login. Two accounts exist already:

- johnsmith (password: 123)
- pauljones3 (password: 123)

Or you can create your own account - on the right-hand side see the *Registration* box. Click the `Create One` link. Enter any details you like and most important: give yourself a few million dollars to play with!



Note

Once you have an account you can start trading. If you get an error "No instances available for accounts", the Web application hasn't found the `accounts` microservice.

16.3. Deploying to Cloud Foundry

1. There is already an application manifest

```
$COURSE_HOME/lab_05/lab_05d/solution/spring-trader-web/manifest.yml:
```

```
timeout: 180
instances: 1
memory: 512M
env:
  SPRING_PROFILES_DEFAULT: cloud
  JAVA_OPTS: -Djava.security.egd=file:///dev/urandom
applications:
- name: webtrader
  random-route: true
  path: target/web-1.0.0-SNAPSHOT.jar
  services: [ service-registry ]
```

However, before we can push to Cloud Foundry we need to tell the application where to find its microservices.

In your IDE or a suitable editor, open `src/main/resources/application.yml`. It looks like this:

```
spring:
  profiles.active: local
  application:
    name: web

---
spring:
  profiles: local

  jpa:
    hibernate:
      ddl-auto: create-drop
server:
  port: 8080

quotes:
  url: http://localhost:8086

accounts:
  url: http://localhost:8082

portfolios:
  url: http://localhost:8081

---
spring:
  profiles: cloud
  jpa:
    hibernate:
      ddl-auto: update

quotes:
  url: http://quotes-xxxx-yyyy.cfapps.io

accounts:
  url: http://accounts-xxxx-yyyy.cfapps.io

portfolios:
  url: http://portfolios-xxxx-yyyy.cfapps.io
```

In the last section (`spring.profiles: cloud`), you need to change the `url` properties for the URLs of the three microservices you previously pushed to Cloud Foundry.

Make sure these applications are still running in Cloud Foundry.

2. Push to Cloud Foundry:

```
$ cf push
...
```

3. Access the application on `http://webtrader<your-random-words>.cfapps.io`.

Chapter 17. Lab 06a - Deploying and Using Spring Cloud Config Server

Estimated time to complete: 30 minutes

Now we'll begin using the components found in Spring Cloud to implement patterns for distributed systems. We'll be using Spring Cloud Config to centralize configuration for applications. Our configuration will be stored in a local folder, and served to our applications using the Spring Cloud Config Server.

In this lab, we'll do the following:

1. Create a folder to store your configuration
2. Create a Spring Cloud Config Server and test it with a basic sample application



Note

before starting this lab, you can stop all the applications you had started in the previous labs. You will not need them anymore.

This lab uses 4 different `properties` files, as described in this table:

File name	Location	Purpose
<code>demo.properties</code>	Config folder	Configuration file for application <code>demo</code>
<code>demo-dev.properties</code>	Config folder	<code>dev</code> profile configuration for application <code>demo</code>
<code>bootstrap.properties</code>	<code>springtrader-config-client</code>	Information sent to config server before application starts
<code>application.properties</code>	all Spring Boot projects	Spring Boot configuration

17.1. Create the config folder to Store Configuration

1. Change to the lab directory:

```
$ cd $COURSE_HOME/lab/lab_06/lab_06a/initial/
```

2. Create a directory for the configuration files

```
mkdir config  
cd config
```

3. Create application.properties

```
greeting=Hello world from application.properties  
val_1=value from application.properties
```

4. Create demo.properties

```
greeting=Hello world from demo.properties  
val_2=value from demo.properties
```

5. Create demo-dev.properties

```
greeting=Hello world from demo-dev.properties  
val_3=value from demo-dev.properties
```

17.2. Create a Spring Cloud Config Server

1. Change to the lab directory:

```
$ cd $COURSE_HOME/lab/lab_06/lab_06a/initial/springtrader-config-server
```

and import the project (via pom.xml) into your IDE of choice.

2. Open pom.xml. As you can see, your project has been configured for Spring Cloud already. Add a dependency on spring-cloud-config-server:

```
<dependency>  
  <groupId>org.springframework.cloud</groupId>  
  <artifactId>spring-cloud-config-server</artifactId>  
</dependency>
```

3. Now open io.springtrader.configserver.SpringtraderConfigServerApplication and add the @EnableConfigServer annotation:

```
@SpringBootApplication  
@EnableConfigServer // <--- ADD THIS!  
public class SpringtraderConfigServerApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(SpringtraderConfigServerApplication.class, args);  
    }  
}
```

4. Now open src/main/resources/application.properties and paste the following code:

```
server.port=9000  
spring.cloud.config.server.native.searchLocations=file:///config/
```

```
spring.profiles.active=native
```

5. Run the application

6. Test the application in your web browser by connecting to <http://localhost:9000/application/default> You should be able to see your properties displayed similar to this:

```
{
  "name": "application",
  "profiles": [
    "default"
  ],
  "label": "master",
  "propertySources": [
    {
      "name": "file../config/application.properties",
      "source": {
        "val_1": "value from application.properties",
        "greeting": "Hello world from application.properties"
      }
    }
  ]
}
```

As shown above, you should be able to see the `greeting` value that you have entered inside `application.properties`.

7. Now experiment with following links, pay attention to the value of `greeting` key

- <http://localhost:9000/application-default.properties>
- <http://localhost:9000/demo-default.properties>
- <http://localhost:9000/demo-dev.properties>



Note

<http://localhost:9000/<myApp>/default> displays properties for all managed applications.

Leave the Config Server running, the client will need it.

17.3. Create the Sample Test Application

1. Change to the lab directory:

```
$ cd $COURSE_HOME/lab/lab_06/lab_06a/initial/springtrader-config-client
```

and import the project (via `pom.xml`) into your IDE of choice.

2. Open `pom.xml` and add the following dependency:

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-config</artifactId>
</dependency>
```



Note

`spring-cloud-starter-config` comes with `spring-cloud-config-client` as a transitive dependency

1. In the package `io.springtrader.configclient` create the class `GreetingController`, and into that class paste the following source code:

```
@RestController  
public class GreetingController {  
  
    @Value("${greeting}")  
    String greeting;  
  
    @RequestMapping("/")  
    public String greeter() {  
        return greeting;  
    }  
}
```

2. Inside `src/main/resources` open the `bootstrap.properties` file and set the config server's uri:

```
spring.cloud.config.uri=http://localhost:9000
```

3. Make sure the config server is still running. Launch the config-client application. Test the client has succeeded in fetching its configuration by using your web browser to connect to <http://localhost:8080>. You should be able to see a "Hello World" message displayed.

```
Hello world from application.properties
```

4. Add the application name to `bootstrap.properties`

```
spring.cloud.config.uri=http://localhost:9000  
spring.application.name=demo
```

5. Restart the client and check that the hello world message has been changed

```
Hello world from demo.properties
```

6. Add the dev profile to `bootstrap.properties`

```
spring.cloud.config.uri=http://localhost:9000  
spring.application.name=demo  
spring.profiles.active=dev
```

7. Restart the client and check that the hello world message has been changed again

```
Hello world from demo-dev.properties
```

8. Take a look at the Spring Environment to see how the `greeting` property is being resolved. You can connect to <http://localhost:8080/env> The information you need is on the first line.

17.4. [Bonus] Use github to store the config files

1. Create a github repository
2. Push the config folder to the github repository (make sure the properties files are at the root of the repository, for example <https://github.com/pivotal-education/springtrader-config-repo>)
3. Change the config server's application.properties to

```
server.port=9000
spring.cloud.config.server.git.uri= https://github.com/<YourUserName>/<RepoName>.git
```
4. Restart the config server
5. Restart the client to see the values are properly loaded
6. Push some changes to github
7. Restart the client again to check the changes

Chapter 18. Lab 06b - Leveraging Eureka for Service Discovery via Spring Cloud Netflix

Estimated time to complete: 30 minutes

Let's continue learning the components found in Spring Cloud to implement patterns for distributed systems. We'll use Spring Cloud Netflix to deploy Eureka, which is a component offering service registration and discovery.

In this lab, we'll do the following:

1. Create a Eureka server
2. Create two applications, a `producer` and a `consumer`, wire them up using Eureka, and test them.

18.1. Introduction

1. Before starting this lab, you can stop all the web applications that you had started in the previous labs. You will not need them anymore.
2. There are four different web applications to be run locally in this lab. To avoid any confusion, we have listed all the port numbers below:

Application	port number
Eureka Server	8761
Producer	8080 (default)
Consumer	8091

18.2. Creating a Eureka Server

1. Change to the lab directory:

```
cd $COURSE_HOME/lab/lab_06/lab_06b/initial/
```

Lab 06b - Leveraging Eureka for Service Discovery via Spring Cloud Netflix

and import all 4 projects into your IDE. For the first part of this lab, we will work on `springtrader-eureka`

2. Open `pom.xml`, change the parent POM to the `spring-cloud-starter-parent`:

```
<parent>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-parent</artifactId>
  <version>Brixton.RELEASE</version>
</parent>
```

3. Add a dependency on `spring-cloud-starter-eureka-server`:

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-eureka-server</artifactId>
</dependency>
```

4. In `io.springtrader.eureka.SpringtraderEurekaApplication`, add the `@EnableEurekaServer` annotation:

```
@SpringBootApplication
@EnableEurekaServer // <-- ADD THIS!
public class SpringtraderEurekaApplication {

    public static void main(String[] args) {
        SpringApplication.run(SpringtraderEurekaApplication.class, args);
    }
}
```

5. Open `src/main/resources/application.yml` and paste in the following source:

```
server:
  port: 8761

eureka:
  instance:
    hostname: localhost
  client:
    registerWithEureka: false
    fetchRegistry: false
    serviceUrl:
      defaultZone: http://${eureka.instance.hostname}:${server.port}/eureka/
```

6. Run the application and visit <http://localhost:8761>

You should see the below console:

The screenshot shows the Spring Cloud Netflix dashboard. In the 'System Status' section, it displays environment information like Data center, current time (2015-02-18T18:49:44 -0800), uptime (00:00), lease expiration enabled (false), renew threshold (0), and renew frequency (0). The 'DS Replicas' section shows a single instance registered under 'localhost'. The 'Instances currently registered with Eureka' table has columns for Application, AMIs, Availability Zones, and Status, with a note that no instances are available.

18.3. Create and Register the Producer Service

1. In your IDE, open the `lab_06b-springtrader-producer` project
2. Open `pom.xml`, change the parent POM to the `spring-cloud-starter-parent`:

```
<parent>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-parent</artifactId>
  <version>Brixton.RELEASE</version>
</parent>
```

3. Add a dependency on `spring-cloud-starter-eureka`:

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-eureka</artifactId>
</dependency>
```

4. In the package `io.springtrader.producer`, create the class `ProducerController`. Into that class paste the following code:

```
@RestController
public class ProducerController {

    private org.slf4j.Logger log = org.slf4j.LoggerFactory.getLogger(ProducerController.class);
    private AtomicInteger counter = new AtomicInteger(0);

    @RequestMapping("counter-value")
    public Object produce() {
        int value = counter.getAndIncrement();
        log.info("Produced a value: {}", value);

        return Collections.singletonMap("value", value);
    }
}
```

Lab 06b - Leveraging Eureka for Service Discovery via Spring Cloud Netflix

5. Now open `io.springtrader.producer.SpringtraderProducerApplication` and add the `@EnableDiscoveryClient` annotation:

```
@SpringBootApplication
@EnableDiscoveryClient // <---- ADD THIS!
public class SpringtraderProducerApplication {

    public static void main(String[] args) {
        SpringApplication.run(SpringtraderProducerApplication.class, args);
    }
}
```

6. Create the file `src/main/resources/application.yml` and paste in the following source:

```
spring:
  application:
    name: producer

server:
  port: 9009

eureka:
  instance:
    preferIpAddress: true
    leaseRenewalIntervalInSeconds: 10
    metadataMap:
      instanceId: ${server.port}
  client:
    serviceUrl:
      defaultZone: http://localhost:8761/eureka/
```

7. Run the producer application

8. Test your producer application <http://localhost:9009/counter-value> refresh several times to see the counter value moves up

9. Ten seconds after the producer application finishes startup, you should see it log its registration with Eureka:

```
DiscoveryClient_PRODUCER/192.168.1.10:9009 - Re-registering apps/PRODUCER
DiscoveryClient_PRODUCER/192.168.1.10:9009: registering service...
DiscoveryClient_PRODUCER/192.168.1.10:9009 - registration status: 204
```

You should also be able to refresh <http://localhost:8761> in the browser and see the producer registered:

Instances currently registered with Eureka			
Application	AMIs	Availability Zones	Status
PRODUCER	n/a (1)	(1)	UP (1) - potiguar.local:producer:9009

18.4. Create and Register the Consumer Service

1. In your IDE, open `lab_06b-springtrader-consumer`
2. Open its `pom.xml` file and change the parent POM to `spring-cloud-starter-parent`:

```
<parent>
  <groupId>org.springframework.cloud</groupId>
```

Lab 06b - Leveraging Eureka for Service Discovery via Spring Cloud Netflix

```
<artifactId>spring-cloud-starter-parent</artifactId>
<version>Brixton.RELEASE</version>
</parent>
```

3. Add a dependency on spring-cloud-starter-eureka:

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-eureka</artifactId>
</dependency>
```

4. In the package io.springtrader.consumer, create the class ProducerResponse, and into that class paste the following code:

```
public class ProducerResponse {
    private int value;

    public void setValue(int value) {
        this.value = value;
    }

    public int getValue() {
        return value;
    }
}
```

5. Also in the package io.springtrader.consumer.controller, create the class ConsumerController, and into that class paste the following code:

```
@RestController
public class ConsumerController {

    @Autowired
    DiscoveryClient discoveryClient;

    RestTemplate restTemplate = new RestTemplate();

    @RequestMapping("/")
    Object consume2() {
        ServiceInstance instance = discoveryClient.getInstances("PRODUCER").get(0);

        ProducerResponse response = restTemplate.getForObject(
            instance.getUri() + "/counter-value", ProducerResponse.class);

        return Collections.singletonMap("value", response.getValue());
    }
}
```

6. Now open io.springtrader.consumer.SpringtraderConsumerApplication and add the @EnableDiscoveryClient annotation:

```
@SpringBootApplication
@EnableDiscoveryClient // <--- ADD THIS!
public class SpringtraderConsumerApplication {

    public static void main(String[] args) {
        SpringApplication.run(SpringtraderConsumerApplication.class, args);
    }
}
```

7. In src/main/resources/application.yml, set the server.port property:

```
spring:
  application:
    name: consumer

server:
  port: 8091
```

Lab 06b - Leveraging Eureka for Service Discovery via Spring Cloud Netflix

```
eureka:  
  instance:  
    preferIpAddress: true  
    leaseRenewalIntervalInSeconds: 10  
    metadataMap:  
      instanceId: ${server.port}  
  client:  
    serviceUrl:  
      defaultZone: http://localhost:8761/eureka/
```

8. Run the consumer application

9. Ten seconds after the consumer application finishes startup, you should see it log its registration with Eureka:

```
DiscoveryClient_CONSUMER/192.168.1.10:8091 - Re-registering apps/CONSUMER  
DiscoveryClient_CONSUMER/192.168.1.10:8091: registering service...  
DiscoveryClient_CONSUMER/192.168.1.10:8091 - registration status: 204
```

You should also be able to refresh <http://localhost:8761> in the browser and see the producer registered:

Instances currently registered with Eureka

Application	AMIs	Availability Zones	Status
CONSUMER	n/a (1)	(1)	UP (1) - potiguar.local:consumer:8091
PRODUCER	n/a (1)	(1)	UP (1) - potiguar.local:producer:9009

10. Open a browser tab on the consumer application (<http://localhost:8091>). It should show that it is receiving values from the producer:

```
{  
  "value": 0  
}
```

Chapter 19. Lab 07a - Client-Side Load Balancing with Ribbon

Estimated time to complete: 30 minutes

Let's continue learning the components found in Spring Cloud to implement patterns for distributed systems. We'll again use Spring Cloud Netflix to implement client-side load balancing with Ribbon.

In this lab, we'll do the following:

1. Change the consumer application from [discovery lab](#) to:
 - a. Use a `LoadBalancerClient`
 - b. Use a `RestTemplate` configured to resolve service names from Ribbon
2. Test the new consumer versions against our local pool of producers
3. Try swapping out the Ribbon load balancing algorithm



Note

before starting this lab, you can stop all the applications you had started in the previous labs. You will not need them anymore.

This lab uses 4 running applications. Here is a table that you can refer to during the lab.

Eureka	http://localhost:8761
Producer 1	http://localhost:8080
Producer 2	http://localhost:8082
Consumer	http://localhost:8091

19.1. Setup

1. Four projects have been placed inside `$COURSE_HOME/lab_07/lab_07a/initial/`. Import all of them inside your IDE of choice.
2. Eureka: Open `springtrader-eureka` and launch the class `SpringtraderEurekaApplication`
3. Connect to <http://localhost:8761> to confirm that Eureka has been started successfully.
4. A completed `springtrader-producer` project has been placed in `$COURSE_HOME/lab_07/lab_07a/initial/springtrader-producer` for your convenience. Change to that directory and build

```
cd $COURSE_HOME/lab/lab_07/lab_07a/initial/springtrader-producer  
mvn clean package
```

5. Now run the producer twice (on different ports), each in a different terminal/command window:

```
Linux, Mac OS:  
$ SERVER_PORT=8080 java -jar target/lab_07a-springtrader-producer-0.0.1-SNAPSHOT.jar  
$ SERVER_PORT=8082 java -Deureka.instance.metadataMap.instanceId=p8082 -jar target/lab_07a-springtrader-producer-0.0.1-SNAPSHOT.jar  
  
Windows:  
> set SERVER_PORT=8080  
> java -jar target/lab_07a-springtrader-producer-0.0.1-SNAPSHOT.jar  
> set SERVER_PORT=8082  
> java -Deureka.instance.metadataMap.instanceId=p8082 -jar target/lab_07a-springtrader-producer-0.0.1-SNAPSHOT.jar
```

6. Ensure you have two instances of the producer service registered in Eureka on <http://localhost:8761>:

Application	AMIs	Availability Zones	Status
PRODUCER	n/a (2)	(2)	 UP (2) - turkey:producer:8080 , turkey:producer:8082

19.2. Using the LoadBalancerClient

1. Inside your IDE, open the project `springtrader-consumer`
2. Where is Ribbon?

You might be interested to know how your Ribbon Client has been included inside your project so it can be used by our Consumer. Inside your `pom.xml` file, you will find the following dependency:

```
<dependency>  
  <groupId>org.springframework.cloud</groupId>  
  <artifactId>spring-cloud-starter-eureka</artifactId>  
</dependency>
```

The `spring-cloud-starter-eureka` dependency transitively comes with the Ribbon libraries (you will find them in your list of jar files).

1. In `io.springtrader.consumer.ConsumerController`, include the `consume()` method as follows:

```
@RequestMapping(value = "/", produces = "application/json")
public String consume() {
    ProducerResponse response = restTemplate.getForObject("http://producer", ProducerResponse.class);
    return String.format("%s\\server port\\%d, %s\\value\\%d", response.getServerPort(), response.getValue());
}
```

2. Config RestTemplate, add following codes to `SpringtraderConsumerApplication`

```
@LoadBalanced
@Bean
RestTemplate restTemplate() {
    return new RestTemplate();
}
```

3. Run your application by launching `SpringtraderConsumerApplication`.

4. Connect to <http://localhost:8091> and access the consumer application several times. Per the round-robin algorithm, 50% of the requests should be served by producer 8080 (and the other half should go to producer 8082).

```
{
    "server port":8082,
    "value":4
}
```

19.3. Failing one of the producers

You now have a Load Balancer that works across 2 Producers. But what happens when one of the Producers fails?

1. Let's simulate a malfunctioning producer. Open the link <http://localhost:8080/setError/true>, this will make the <http://localhost:8080> always returns an error.
2. Now try to access the consumer a few times. What can you notice?
3. In the next lab, you will learn how to use a `circuit breaker` so failures can be handled gracefully.

Chapter 20. Lab 07b - Fault-Tolerance with Hystrix

Estimated time to complete: 20 minutes

Let's continue learning the components found in Spring Cloud to implement patterns for distributed systems. We'll again use Spring Cloud Netflix to implement client-side load balancing with Ribbon. However, this time we'll add an implementation of the Circuit Breaker pattern using Netflix Hystrix.

In this lab, we'll do the following:

1. Refactor the consumer application from the [Load Balancing](#) lab to add a `ProducerClient`.
2. Test the new consumer version against our local pool of producers
3. Kill one of the producers and watch the consumer failover.
4. Set the remaining producer into the error state, and watch the circuit breaker fallback behavior.
5. Kill the remaining producer and watch the circuit breaker fallback behavior.
6. Restart a producer and watch the consumer recover.

20.1. Setup

1. You can just reuse the Eureka server and the 2 producers from Lab 07a as shown below:

Eureka	http://localhost:8761
Producer 1	http://localhost:8080
Producer 2	http://localhost:8082

In case you did not have time to complete the lab 07a, you can run all 3 application instances from lab 07a's `complete` folder.

1. Ensure you have two instances of the producer service registered in Eureka (<http://localhost:8761>):

Application	AMIs	Availability Zones	Status
PRODUCER	n/a (2)	(2)	UP (2) - turkey:producer:8080 , turkey:producer:8082

20.2. Using Hystrix



Note

all the failover configuration is done on the consumer side (as opposed to the producer side. In case the producer fails, the consumer should know where to fall back to.

1. Inside your IDE, go into the `springtrader-consumer` project.

2. Open its `pom.xml` file and add the following dependency:

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-hystrix</artifactId>
</dependency>
```

3. You now need to add the relevant configuration for your Circuit Breaker. We will rely on Spring Boot's default configuration.

Inside the class `io.springtrader.consumer.SpringtraderConsumerApplication`, add the annotation `@EnableCircuitBreaker`.

4. Config RestTemplate, add following codes to `SpringtraderConsumerApplication`

```
@LoadBalanced
@Bean
RestTemplate restTemplate() {
    return new RestTemplate();
}
```

5. Create the class `io.springtrader.consumer.ProducerClient` and into it paste the following code:

```
@Component
public class ProducerClient {

    @Autowired
    @LoadBalanced
    RestTemplate restTemplate;

    @HystrixCommand(fallbackMethod = "getProducerFallback")
    public ProducerResponse getValue() {
        return restTemplate.getForObject("http://producer", ProducerResponse.class);
    }

    public ProducerResponse getProducerFallback() {
        ProducerResponse response = new ProducerResponse();
        response.setValue(-1);
        response.setServerPort(-1);
        return response;
    }
}
```

6. Refactor the class `io.springtrader.consumer.ConsumerController` to autowire the `ProducerClient` instead of `RestTemplate`, and then use it to obtain the `ProducerResponse`:

```
@Autowired  
ProducerClient client;  
  
@RequestMapping(value = "/*", produces = "application/json")  
String consume() {  
    ProducerResponse response = client.getValue();  
    return String.format("{\"value\":%d}", response.getValue());  
}
```

7. Inside `src/main/resources/application.yml` add the following:

```
producer.ribbon.ServerListRefreshInterval: 5000
```

8. Start the application by launching `SpringtraderConsumerApplication`

In your logs, you should see that Hystrix has been loaded:

```
c.n.h.c.m.e.HystrixMetricsPoller : Starting HystrixMetricsPoller  
ration$HystrixMetricsPollerConfiguration : Starting poller
```

1. Test the consumer application on <http://localhost:8091> and show that it is still receiving values from the producers. Also, watch the producers and observe that Ribbon is performing round robin load balancing across the two producers

20.3. Failover

1. Shut down one of the two producer processes.
2. Reload url, test the consumer application and show that it is still receiving values from one of the producers. You may get the fallback value of `-1` a few times due to the lag in removing the failed instance from the Ribbon cache. Eventually it will converge to hitting the only remaining healthy instance.

20.4. Fallback

1. Send the remaining producer into the error state by visiting `/setError/true`, now access the consumer, you should **always** see the fallback value `-1`.
2. Shut down the remaining producer process.
3. Reload url, test the consumer application and show that it is only emitting the fallback value of `-1`.

20.5. Recovery

1. Restart one of the producer processes. Wait for it to register with Eureka.
2. Reload url, test the consumer application and show that eventually recovers and starts hitting the new producer process. This can take several seconds as the Eureka and Ribbon caches repopulate.

20.6. Next Steps

Do not shut anything down when you complete this lab. We will add one additional component in [next lab](#)

Chapter 21. Lab 07c - Monitoring Circuit Breakers with Hystrix Dashboard

Estimated time to complete: 20 minutes

We will now monitor what's happening in our Circuit Breakers by using Netflix Hystrix Dashboard.

In this lab, we'll do the following:

1. Run everything that we created in the previous lab.
2. Create a Hystrix Dashboard
3. Kill both of the producers and watch the circuit open in the Dashboard.
4. Restart a producer and watch the circuit close in the Dashboard.

21.1. Setup

1. Make sure you've completed the previous lab and have everything still running.

21.2. Building the Hystrix Dashboard

1. The project `springtrader-hystrix-dashboard` has been placed inside `$COURSE_HOME/lab_07/lab_07c/initial/`. Import it inside your IDE of choice.

2. Add the following dependency to `pom.xml`:

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-hystrix-dashboard</artifactId>
</dependency>
```

3. Add a `@EnableHystrixDashboard` to the class `io.pivotal.springcloud.hystrixdashboard.HystrixDashboardApplication`.

4. Paste the following inside `src/main/resources/application.yml`:

```
server:
  port: 7979
```

Lab 07c - Monitoring Circuit Breakers with Hystrix Dashboard

5. Run the application by launching `HystrixDashboardApplication`
6. Now, access the Hystrix Dashboard at <http://localhost:7979/hystrix>. Enter `http://localhost:8091/hystrix.stream` into the text field and click the "Monitor Stream."



Hystrix Dashboard

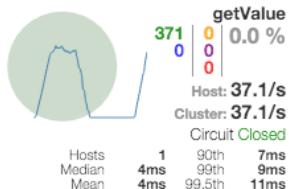
*Cluster via Turbine (default cluster): http://turbine-hostname:port/turbine.stream
Cluster via Turbine (custom cluster): http://turbine-hostname:port/turbine.stream?cluster=[clusterName]
Single Hystrix App: http://hystrix-app:port/hystrix.stream*

Delay: ms Title:

7. Access the consumer application at <http://localhost:8091/> several times and check that the circuit breaker is registering successful requests.

Hystrix Stream: <http://localhost:8081/hystrix.stream>

Circuit Sort: [Error then Volume](#) | [Alphabetical](#) | [Volume](#) | [Error](#) | [Mean](#) | [Median](#) | [90](#) | [99](#) | [99.5](#)



Thread Pools Sort: [Alphabetical](#) | [Volume](#) |

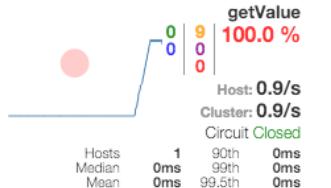


21.3. Fallback

1. Shut down both of the producer processes.
2. Test the consumer application and show that it is only emitting the fallback value of -1.
3. You should also see the circuit breaker is registering short-circuited requests.

Hystrix Stream: <http://localhost:8081/hystrix.stream>

Circuit Sort: [Error then Volume](#) | [Alphabetical](#) | [Volume](#) | [Error](#) | [Mean](#) | [Median](#) | [90](#) | [99](#) | [99.5](#)



Thread Pools Sort: [Alphabetical](#) | [Volume](#) |



21.4. Recovery

1. Restart one of the producer processes. Wait for it to register with Eureka.
2. Test the consumer application and show that eventually recovers and starts hitting the new producer process. This can take up to one minute as the Eureka and Ribbon caches repopulate.
3. Continue to watch the dashboard. You should also see the circuit breaker dashboard show the recovery.

Chapter 22. Lab 07d - Declarative REST Clients with Feign

Estimated time to complete: 15 minutes

This lab is based on the [Load Balancing](#) lab. We will enhance it to use a declarative REST repository based on Feign.

We will do the following:

1. Improve the consumer application from the Load Balancing lab to use a `ProducerClient` enabled by Feign
2. Test the new consumer version against our local pool of producers

22.1. Setup

1. Stop all the running applications and remove the corresponding projects from your IDE
2. Inside your IDE, import the projects that have been placed for you inside `$COURSE_HOME/lab_07/lab_07d/initial/`



Note

those projects are identical to the ones that you had completed at the end of the Lab 07a (Load Balancing with Ribbon).

1. Start your Eureka server and 2 producers as we have done in the previous labs. You SHOULD NOT start the consumer as of now. The below applications should be running:
2. Once you are ready to go, ensure you have two instances of the producer service registered in Eureka (<http://localhost:8761>):

Application	AMIs	Availability Zones	Status
PRODUCER	n/a (2)	(2)	UP (2) - turkey:producer:8080 , turkey:producer:8082

22.2. Using Feign

1. Open your `springtrader-consumer` project inside your IDE. Add the following dependency to `pom.xml`:

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-feign</artifactId>
</dependency>
```

2. Add a `@EnableFeignClients` annotation to the `SpringtraderConsumerApplication`.

3. Create the interface `io.springtrader.consumer.ProducerClient` and into it paste the following code:

```
@FeignClient("producer")
public interface ProducerClient {
    @RequestMapping(method = RequestMethod.GET, value = "/*")
    ProducerResponse getValue();
}
```

4. Refactor the class `io.springtrader.consumer.ConsumerController` to autowire the `ProducerClient` instead of `RestTemplate`, and then use it to obtain the `ProducerResponse`:

```
@Autowired
ProducerClient client;

@RequestMapping(value="/", produces="application/json")
public String consume() {

    ProducerResponse response = client.getValue();
    return String.format("{\"server port\":%d, \"value\":%d}", response.getServerPort(), response.getValue());
}
```

5. Start the `springtrader-consumer` application by launching `SpringtraderConsumerApplication`

6. Test the consumer application on <http://localhost:8091/> and show that it is still receiving values from the producers. It should behave exactly like in the lab 07a. The only difference is that your REST binding code is slightly more concise and elegant.

Chapter 23. Lab 08a - Creating an OAuth2 Authorization Server

[abstract]

We will now create an OAuth 2.0 authorization server. Using spring cloud security support we will see how easy it is to create a working authorization server with very little code needed

23.1. Default behaviour

1. The project `lab_08a-springtrader-auth-server` has been placed inside `$COURSE_HOME/lab_08/lab_08a/initial/`. Import it inside your IDE of choice.
2. Run the application by running `SpringtraderAuthServerApplication`

See what happens before you setup your authentication server. Run the following command:

```
curl http://localhost:8000/auth-server/oauth/token \
-u acme:acmesecret \
-d grant_type=password \
-d username=mstine \
-d password=secret
```

You should see the following output:

```
{
  "timestamp":1463732463298,
  "status":405,
  "error":"Method Not Allowed",
  "exception":"org.springframework.web.HttpRequestMethodNotSupportedException",
  "message":"Request method 'POST' not supported",
  "path":"/auth-server/oauth/token"
}
```

23.2. Configuring the dependencies

1. Add the following dependencies to your pom.xml file

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-oauth2</artifactId>
</dependency>
```

23.3. Enabling the auth server

1. Update SpringtraderAuthServerApplication as shown below:

```
@SpringBootApplication
@EnableAuthorizationServer //// (1)
public class SpringtraderAuthServerApplication {

    public static void main(String[] args) {
        SpringApplication.run(SpringtraderAuthServerApplication.class, args);
    }

    @Configuration
    protected static class OAuth2Config extends AuthorizationServerConfigurerAdapter {

        @Autowired
        private AuthenticationManager authenticationManager;

        @Override
        public void configure(AuthorizationServerSecurityConfigurer oauthServer) throws Exception {
            oauthServer.checkTokenAccess("permitAll()"); //// (4)
        }

        @Override
        public void configure(AuthorizationServerEndpointsConfigurer endpoints) throws Exception {
            endpoints.authenticationManager(authenticationManager);
        }

        @Override
        public void configure(ClientDetailsServiceConfigurer clients) throws Exception {
            clients.inMemory()
                .withClient("acme")
                .secret("acmesecret")
                .authorizedGrantTypes("authorization_code", "refresh_token",
                    "password").scopes("openid"); //// (2)
        }
    }

    @Configuration
    protected static class AuthenticationManagerConfiguration extends GlobalAuthenticationConfigurerAdapter {

        @Override
        public void init(AuthenticationManagerBuilder auth) throws Exception {
            auth.inMemoryAuthentication()
                .withUser("mstine").password("secret").roles("USER", "ADMIN").and()
                .withUser("littleidea").password("secret").roles("USER", "ADMIN").and()
                .withUser("starbuxman").password("secret").roles("USER", "ADMIN"); //// (3)
        }
    }
}
```

1. This bootstraps the auth server by creating Spring filters to handle the token credentials that are passed over requests
2. Register a single client
3. Register several users, other Authorization managers should be used such as JDBC or LDAP based here
4. enable /oauth/check_token

23.4. Testing your Oauth server

1. Restart your application

To test the authentication server execute again a simple POST to the `/auth-server/oauth/token` endpoint, notice that `1. acme:acmesecret` is the client's credential `1. mstine:secret` is the user's credential (Resource Owner in OAuth2 terms)

```
curl http://localhost:8000/auth-server/oauth/token \
-u acme:acmesecret \
-d grant_type=password \
-d username=mstine \
-d password=secret
```

You should get back an authentication token

```
{
  "access_token": "661aac97-55ca-49a0-b8b6-a4a1d8cb63de",
  "token_type": "bearer",
  "refresh_token": "9a605803-4013-4818-ae24-22de7b399018",
  "expires_in": 43199,
  "scope": "openid"
}
```

Now let's do token verification against the authorization server. You can replace <your_token_id> with your value for access_token:

```
curl http://localhost:8000/auth-server/oauth/check_token?token=<your_token_id>
```

You should expecting to see the below:

```
{
  "exp": 1463675798,
  "user_name": "mstine",
  "authorities": [
    "ROLE_ADMIN",
    "ROLE_USER"
  ],
  "client_id": "acme",
  "scope": [
    "openid"
  ]
}
```

You're all set! In the next lab, you will use the same mechanism to have the Resource Server verify your token.

Chapter 24. Lab 08b - Securing a Resource Server with Spring Cloud Security

Now that we have an authorization server from lab 08a, we will create a simple restful webservice and protect it using oauth tokens.

24.1. Running it

1. Import the project `springtrader-resource-server` inside your IDE
2. Run the corresponding application by launching `SpringtraderResourceServerApplication`
3. Access the endpoint on your browser <http://localhost:8001/greeting> and you should receive back a simple json output

```
{  
    "content": "Hello user"  
}
```

24.2. Securing it

1. Add the following dependencies to your pom.xml

```
<dependency>  
    <groupId>org.springframework.cloud</groupId>  
    <artifactId>spring-cloud-starter-oauth2</artifactId>  
</dependency>
```

2. Modify you `application.yml`, it should look like this:

```
server:  
  port: 8001  
  
security:  
  oauth2:  
    resource:  
      token-info-uri: http://localhost:8000/auth-server/oauth/check_token  
    client:  
      client-id: acme  
      client-secret: acmesecret
```



Note

Recall that this is the url we use for the last step of the previous lab

3. Modify the `SpringtraderResourceServerApplication` class to enable OAuth2

```
@SpringBootApplication
@EnableResourceServer // <- Add this
public class SpringtraderResourceServerApplication {

    public static void main(String[] args) {
        SpringApplication.run(SpringtraderResourceServerApplication.class);
    }
}
```

`@EnableResourceServer` will enable the proper Spring Security Filters to look for a token on the headers and validate that against the `token-info-url` to obtain the `SecurityPrincipal`

4. Restart your Resource Server

5. Access the endpoint again at <http://localhost:8001/greeting>. You should see the following message showing an unauthorized response:

```
<oauth>
<error_description>
    Full authentication is required to access this resource
</error_description>
<error>unauthorized</error>
</oauth>
```

24.3. Obtain a valid token



Note

the rest of the lab uses `curl` because you will need to customize your request headers. If you have a browser plugin for that, you can choose to do without `curl`.

In order to access the resource now you will need a valid token from your Authorization Server. To do that execute the command:

```
curl http://acme:acmesecret@localhost:8000/auth-server/oauth/token \
-d grant_type=password \
-d username=mstine \
-d password=secret
```

The result should be something like this:

```
{
    "access_token": "5d04666c-11c2-4a4c-9dae-f0def0acc9c1",
    "token_type": "bearer",
    "refresh_token": "361fa725-735f-4853-bb0f-5a8956a01b33",
    "expires_in": 43199,
    "scope": "openid"
```

```
}
```

1. Grab the access_token from that result and now invoke the following command to call your service, replace the Authorization header with the proper token from the previous step:

```
curl -X GET -H "Authorization: bearer 5d04666c-11c2-4a4c-9dae-f0def0acc9c1" http://localhost:8001/greeting
```

This is all that is needed to protect a resource to require OAuth2 tokens to be accessed

Chapter 25. Lab 08c - Consuming a Secured Resource from OAuth2-Aware Client

Now that we have secured a resource in lab 08b, we will create a simple restful oauth2-aware client that will securely invoke the resource.

25.1. Create a OAuth2-Aware Client

1. Import the project lab_08c-springtrader-client to your IDE

2. Add the following dependency to your pom.xml

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-oauth2</artifactId>
</dependency>
```

3. Modify you application.yml, it should look like this:

```
security:
  oauth2:
    client:
      accessTokenUri: http://localhost:8000/auth-server/oauth/token
      userAuthorizationUri: http://localhost:8000/auth-server/oauth/authorize
      clientId: acme
      clientSecret: acmesecret
      scope: openid
      clientAuthenticationScheme: header
    sso:
      login-path: /login
    basic:
      enabled: false
```

4. Modify the ClientApplication class to handle and use the OAuth2 access token:

```
@SpringBootApplication
@EnableOAuth2Client
@RestController ////////////////////////////////////////////////////////////////// (1)
public class ClientApplication {

    @Autowired ////////////////////////////////////////////////////////////////// (2)
    private OAuth2RestOperations restTemplate;

    public static void main(String[] args) {
        SpringApplication.run(ClientApplication.class, args);
    }

    @RequestMapping("/")
    public String home() {
        return restTemplate.getForObject("http://localhost:8001/greeting", String.class);
    }
}
```

1. This annotation enable configuration for an OAuth2 client in a web application that uses Spring Security and wants to use the Authorization Code Grant from one or more OAuth2 Authorization servers.
2. Rest template that is able to make OAuth2-authenticated REST requests with the credentials of the

provided resource.

25.2. Test your client

1. Run your application by launching `ClientApplication`
2. Access the endpoint using a browser at <http://localhost:8080/> and you get back a login screen. Enter 'littleidea' and 'secret'. It will ask you to approve or deny the request. Click approve to view the resource.



Note

littleidea/secret is one of the users that you have created in Lab 08a.

Congratulations! You now have secured your resources and used a client to securely invoke them.

Chapter 26. (Optional) Lab 09a - Deploying the StringTrader Microservices version Application



Note

This lab is designed only for full PCF installation

We will now deploy the SpringTrader application you have been developing. We will use the Web microservices (provided) together with the services you have already built and we will deploy them on Pivotal Cloud Foundry.

26.1. Copying the 3 services

1. Copy the 3 services: Account, Quotes and Portfolio into **initial** folder.

```
cd $COURSE_HOME/lab/lab_09/lab_09a/initial/  
cp -R ../solution/quotes .  
cp -R ../solution/accounts .  
cp -R ../solution/portfolio .
```

26.2. Starting Spring Cloud Services on Pivotal Cloud Foundry

1. Create Eureka service (service-registry):

```
$ cf create-service p-service-registry standard service-registry
```

2. Initialise Eureka service by clicking on the service Manage link.

3. Create Hystrix Dashboard service (circuit-breaker-dashboard):

```
$ cf create-service p-circuit-breaker-dashboard standard circuit-breaker-dashboard
```

4. Initialise Hystrix dashboard by clicking on the service Manage link.

5. Make sure all manifest.yml files define the app memory with 1G.

```
memory: 1G
```

26.3. Deploying the Application

(Optional) Lab 09a - Deploying the StringTrader Microservices version Application

1. Create deploy.sh(bat) file.

```
cf push -f quotes/manifest.yml  
cf push -f accounts/manifest.yml  
cf push -f portfolio/manifest.yml  
cf push -f web/manifest.yml
```

a. If you are using Unix add `#!/bin/sh` at the first file line.

2. Run the deploy command.

```
./deploy.sh (or deploy.bat)
```

```
App started  
  
OK  
  
App webtrader was started using this command `CALCULATED_MEMORY=$(($PWD/.java-buildpack/open_jdk_jre/bin/java-buildpack  
-memory-calculator-2.0.1_RELEASE -memorySizes=metaspace:64m.. -memoryWeights=heap:75,metaspace:10,native:10,stack:5  
-memoryInitials=heap:100%,metaspace:100% -totMemory:$MEMORY_LIMIT) && JAVA_OPTS="-Djava.io.tmpdir=$TMPDIR  
-XX:OnOutOfMemoryError=$PWD/.java-buildpack/open_jdk_jre/bin/killjava.sh $CALCULATED_MEMORY -Djava.security.egd=file:///  
/dev/urandom" && SERVER_PORT=$PORT eval exec $PWD/.java-buildpack/open_jdk_jre/bin/java $JAVA_OPTS -cp $PWD/:$PWD/.java  
-buildpack/spring_auto_reconfiguration/spring_auto_reconfiguration-1.10.0_RELEASE.jar org.springframework.boot.loader.Ja  
rLauncher'  
  
Showing health and status for app webtrader in org instructor / space development as instructor...  
OK  
  
requested state: started  
instances: 1/1  
usage: 1G x 1 instances  
urls: webtrader-nonbiological-bluebonnet.pcf12.cloud.fe.pivotalk.io  
last uploaded: Mon Feb 22 04:33:06 UTC 2016  
stack: cflinuxfs2  
buildpack: java_buildpack_offline  
  
#0 state since cpu memory disk details  
#0 running 2016-02-22 01:34:38 PM 0.0% 700K of 1G 27.0M of 1G
```

26.4. Testing it out

1. Check if all services are registered with Eureka server. Click on Manage service.

(Optional) Lab 09a - Deploying the StringTrader Microservices version Application

APPLICATIONS		Learn More		
STATUS	APP	INSTANCES	MEMORY	
	accounts https://accounts-harus...	1	1GB	»
	portfolio https://portfolio-unwi...	1	1GB	»
	quotes https://quotes-xenolit...	1	1GB	»
	webtrader https://webtrader-misf...	1	1GB	»

SERVICES		Add Service	
SERVICE INSTANCE	SERVICE PLAN	BOUND APPS	
mongodb-dev Documentation Support Delete	MongoDB for PCF Development	1	
config-service Manage Documentation Support Delete	Config Server standard	0	
service-registry Manage Documentation Support Delete	Service Registry standard	4	
mysql-dev Manage Documentation Support Delete	MySQL for Pivotal Cloud Foundry 100 MB Dev	2	
circuit-breaker-dashboard Manage Documentation Support Delete	Circuit Breaker standard	3	

- See all four services registered with Eureka.

(Optional) Lab 09a - Deploying the StringTrader Microservices version Application

The screenshot shows a web-based service registry interface. At the top, there are two navigation links: "Home" (underlined in green) and "History". Below this, the main content area is titled "Service Registry Status". Under this title, there is a section titled "Registered Apps" which contains a table. The table has three columns: "Application", "Availability Zones", and "Status". There are four rows in the table, each representing a registered application: ACCOUNTS, PORTFOLIO, QUOTES, and WEB. All applications are listed as being in the "default (1)" availability zone and are marked as "UP (1)". A red rectangular box highlights the "Registered Apps" section. Below this, another section titled "System Status" contains a table with six parameters and their corresponding values. The parameters are: Current time (value: 2016-02-22T04:12:09 +0000), Lease expiration enabled (value: true), Self-preservation mode enabled (value: false), Renews threshold (value: 0), and Renews in last minute (value: 8).

Application	Availability Zones	Status
ACCOUNTS	default (1)	UP (1)
PORTFOLIO	default (1)	UP (1)
QUOTES	default (1)	UP (1)
WEB	default (1)	UP (1)

Parameter	Value
Current time	2016-02-22T04:12:09 +0000
Lease expiration enabled	true
Self-preservation mode enabled	false
Renews threshold	0
Renews in last minute	8

3. In browser, go to the **webtrader** URL shown during cf push

(Optional) Lab 09a - Deploying the StringTrader Microservices version Application

The screenshot shows the SpringTrader application interface. At the top, there is a navigation bar with links for home, portfolio, trade, and CX. Below the navigation bar, there are two main sections: "IT Vendors" and "FS Organisations". Each section contains a table with columns for Symbol, Price, Change, High, and Low. The "IT Vendors" section shows data for AMD, ORCL, and IBM. The "FS Organisations" section shows data for BAC, WFC, and JPM. Below these sections are two more panels: "Portfolio Summary" and "User Statistics". The "Portfolio Summary" panel displays the number of holdings (2), current value of held shares (\$40,741.00), total purchase value (\$40,741.00), total sold value (.00), total gain/loss (-.00), and total cash (\$905,723.00). The "User Statistics" panel lists the full name (cx), user ID (cx / cx), creation date (Mon Feb 22 04:40:00 UTC 2016), total logins (1), session creation date (Mon Feb 22 04:41:00 UTC 2016), and open balance (\$1,000,000.00).

4. Finally, you can check if Hystrix is working. Checking its dashboard.

The screenshot shows the Hystrix dashboard under the path "instructor > development > circuit-breaker-dashboard". The dashboard has two main sections: "Circuit" and "Thread Pools". The "Circuit" section displays two circuit breakers: "web.getCompanies" and "web.getQuote". Both circuits are currently closed. The "web.getCompanies" circuit has 1 success, 0 errors, and 0 rejections. The "web.getQuote" circuit also has 1 success, 0 errors, and 0 rejections. The "Thread Pools" section shows the "QuotesIntegrationService" thread pool with 8 active threads, 0 blocked threads, and 0 pending tasks. The host and cluster metrics are all at 0.0/s.

Appendix A. Eclipse Tips

A.1. Introduction

This section will give you some useful hints for using Eclipse.

A.2. Package Explorer View

Eclipse's Package Explorer view offers two ways of displaying packages. Flat view, used by default, lists each package at the same level, even if it is a subpackage of another. Hierarchical view, however, will display subpackages nested within one another, making it easy to hide an entire package hierarchy. You can switch between hierarchical and flat views by selecting the menu inside the package view (represented as a triangle), selecting either Flat or Hierarchical from the Package Presentation submenu.

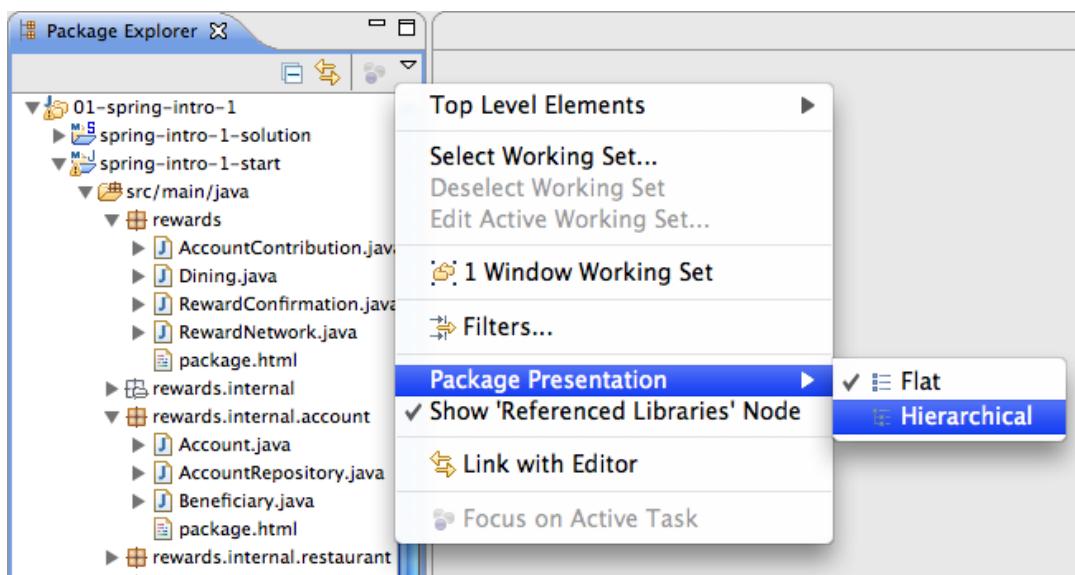


Figure A.1. Switching Views

Switch between hierarchical and flat views by selecting the menu inside the package view (represented as a triangle).

triangle), selecting either Flat or Hierarchical from the Package Presentation submenu

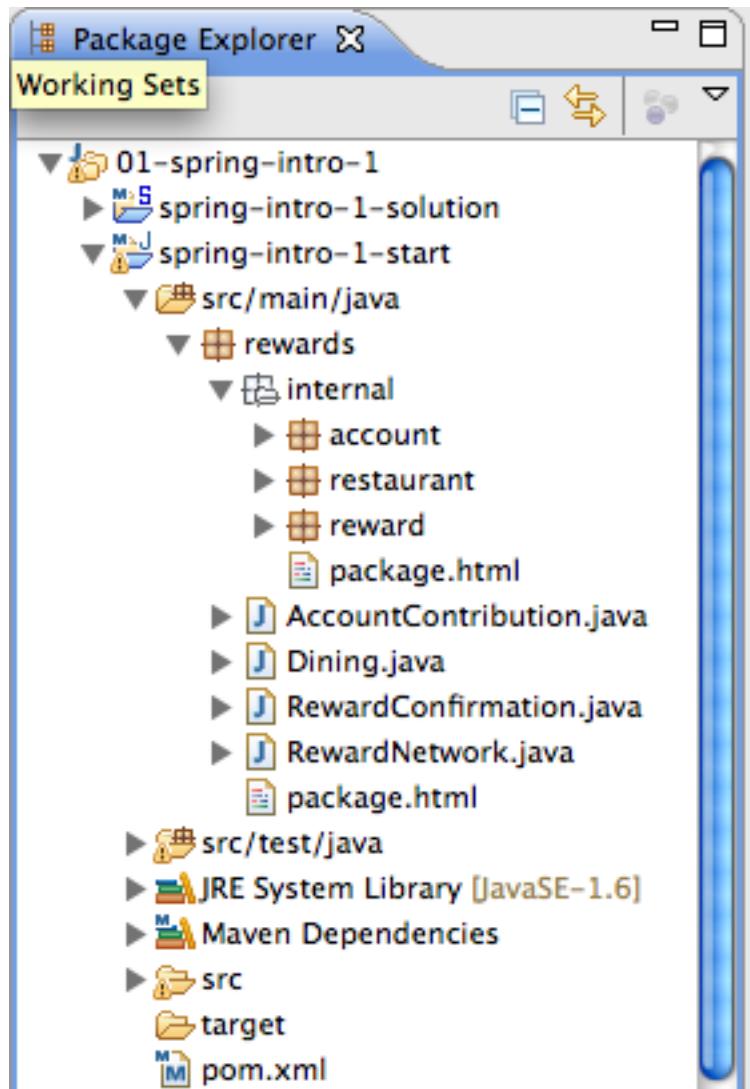


Figure A.2. The hierarchical view shows nested packages in a tree view

A.3. Add Unimplemented Methods

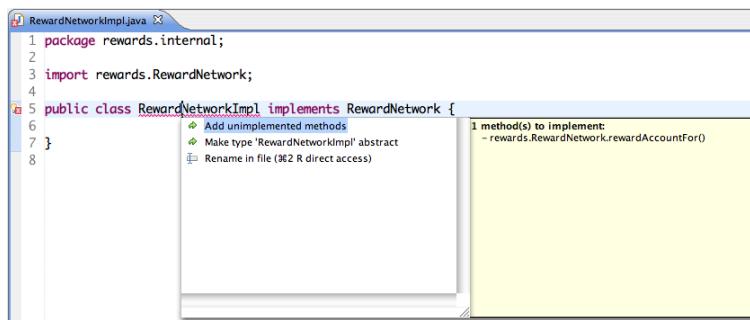


Figure A.3. "Add unimplemented methods" quick fix

A.4. Field Auto-Completion

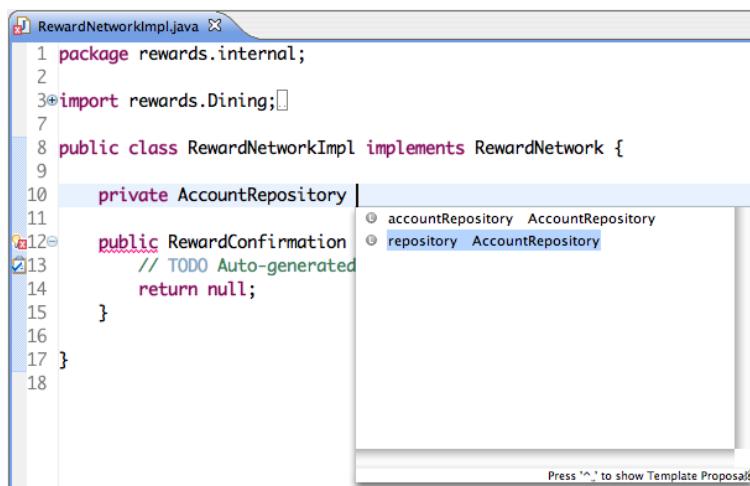


Figure A.4. Field name auto-completion

A.5. Generating Constructors From Fields

You can "Generate a Constructor using Fields" using the Source Menu (ALT + SHIFT + S)

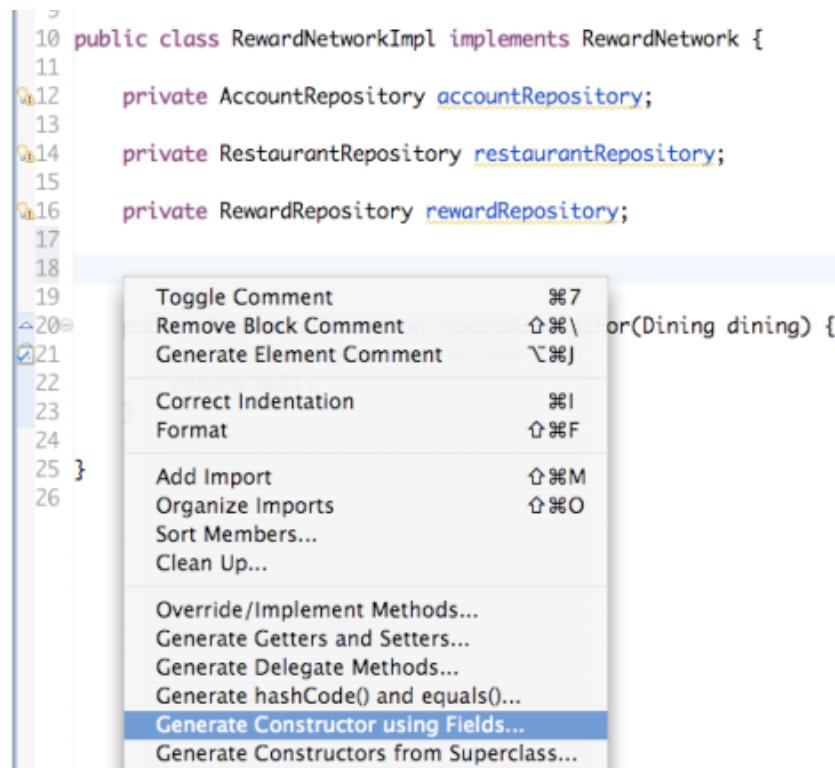
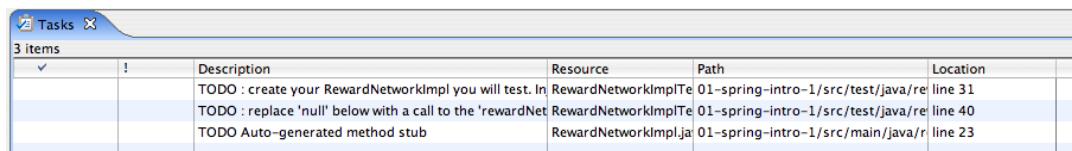


Figure A.5. Generating Constructors

A.6. Field Naming Conventions

A field's name should describe the role it provides callers, and often corresponds to the field's type. It should not describe implementation details. For this reason, a bean's name often corresponds to its service interface. For example, the class JdbcAccountRepository implements the AccountRepository interface. This interface is what callers work with. By convention, then, the bean name should be accountRepository.

A.7. Tasks View



Tasks					
		Description	Resource	Path	Location
	!	TODO : create your RewardNetworkImpl you will test. In	RewardNetworkImplTe	01-spring-intro-1/src/test/java/re	line 31
		TODO : replace 'null' below with a call to the 'rewardNet	RewardNetworkImplTe	01-spring-intro-1/src/test/java/re	line 40
		TODO Auto-generated method stub	RewardNetworkImpl.java	01-spring-intro-1/src/main/java/r	line 23

Figure A.6. The tasks view in the bottom right page area

You can configure the Tasks View to only show the tasks relevant to the current project. In order to do this, open the dropdown menu in the upper-right corner of the tasks view (indicated by the little triangle) and select 'Configure Content...'. Now select the TODO configuration and from the Scopes, select 'On any element in same project'. Now if you have multiple project opened, with different TODOs, you will only see those relevant to the current project.

A.8. Rename a File

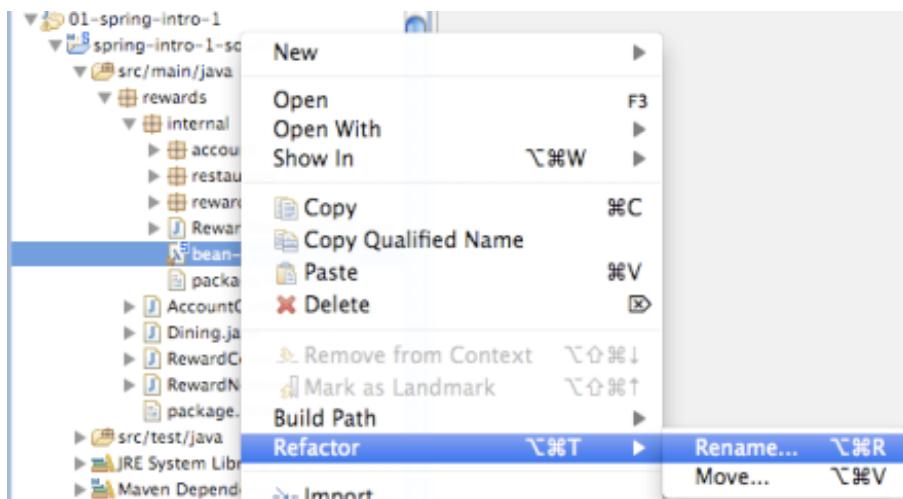


Figure A.7. Renaming a Spring configuration file using the Refactor command